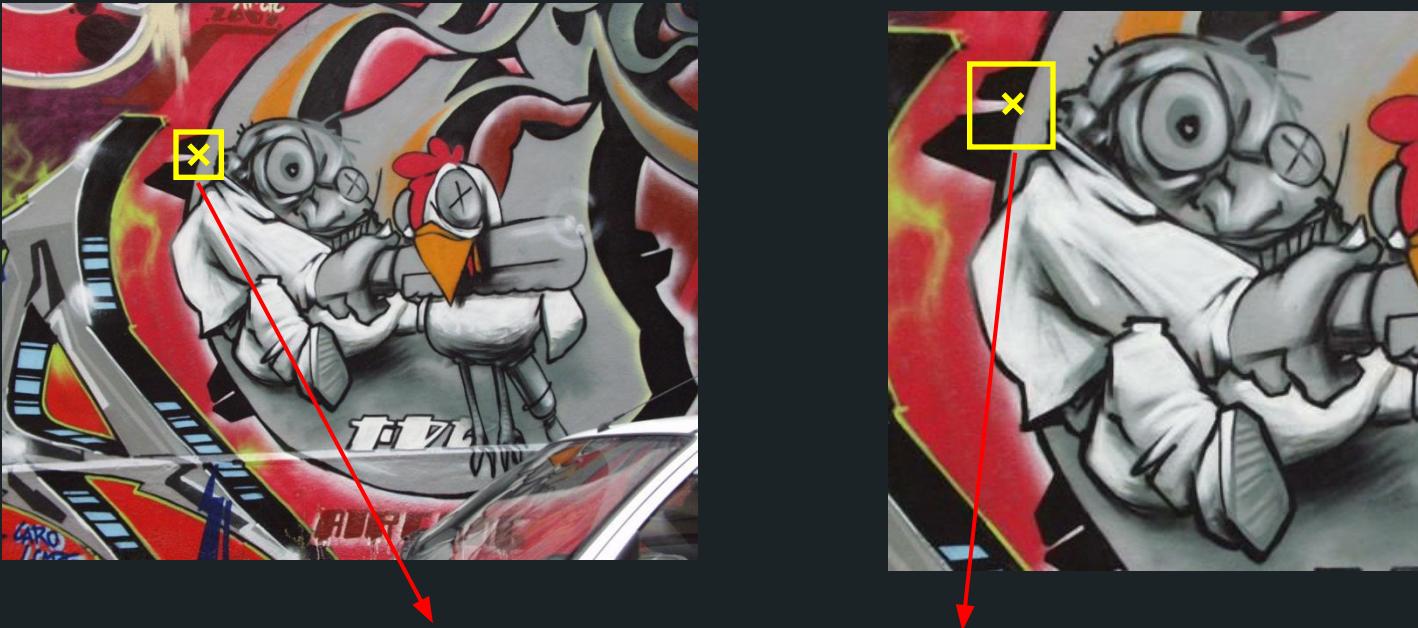


# LOCAL IMAGE FEATURES

So far: can localize in x-y, but not scale



# Automatic Scale Selection

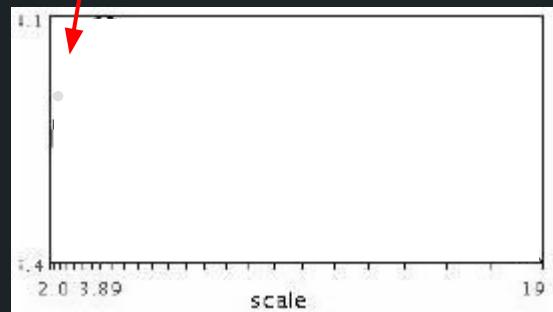


$$f(I_{i_1 \sqcup i_m}(x, \sigma)) = f(I_{i_1 \sqcup i_m}(x', \sigma'))$$

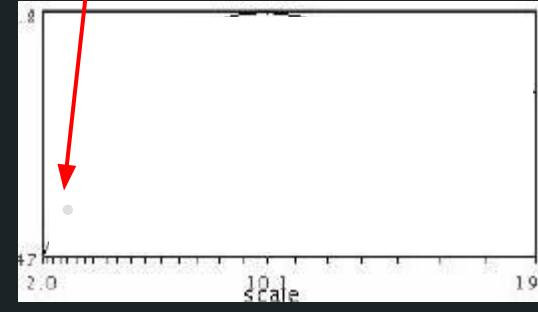
How to find corresponding patch sizes?

# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



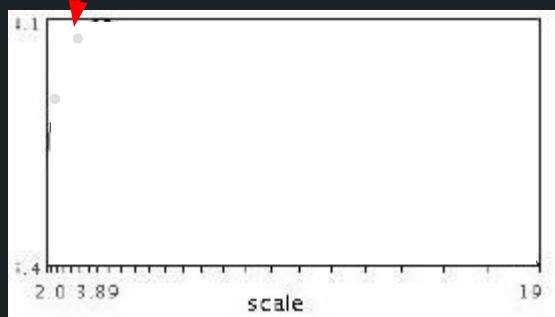
$$f(I_{i \sqcup i_m}(x, \sigma))$$



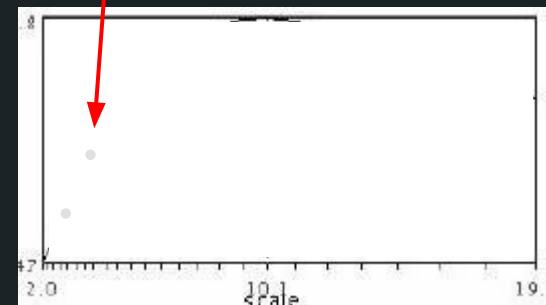
$$f(I_{i \sqcup i_m}(x', \sigma))$$

# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



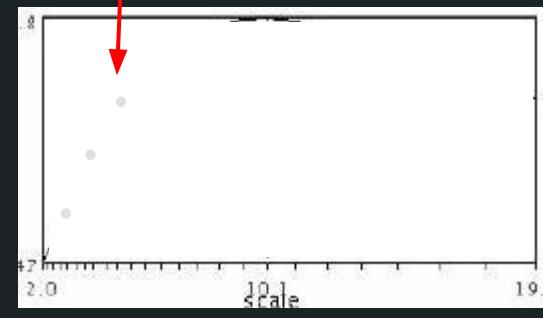
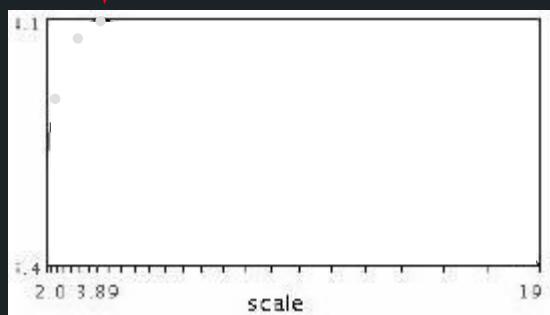
$$f(I_{i \sqcup i_m}(x, \sigma))$$



$$f(I_{i \sqcup i_m}(x', \sigma))$$

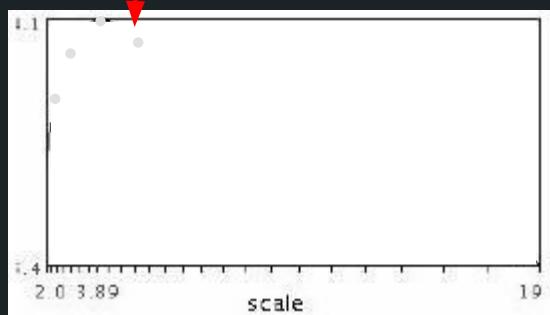
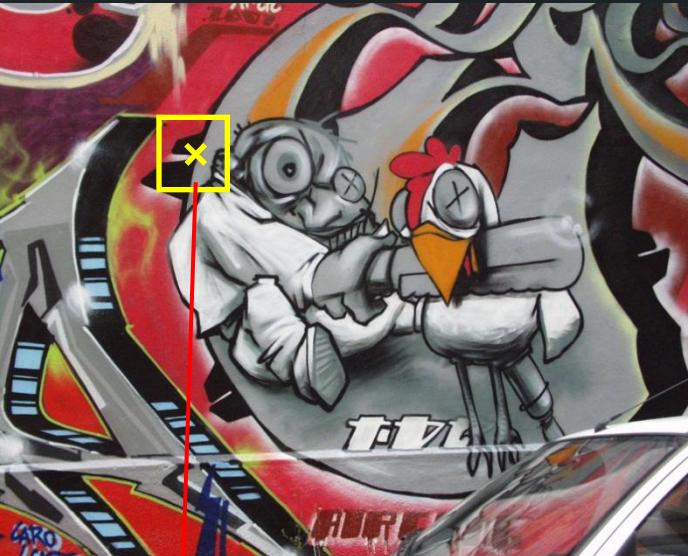
# Automatic Scale Selection

- Function responses for increasing scale (scale signature)

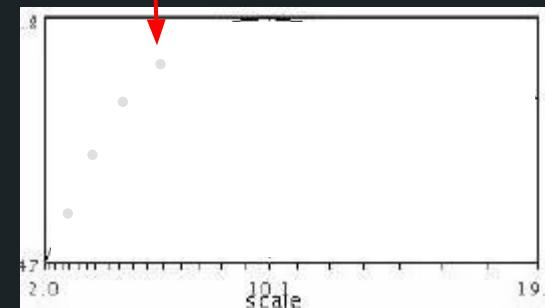


# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



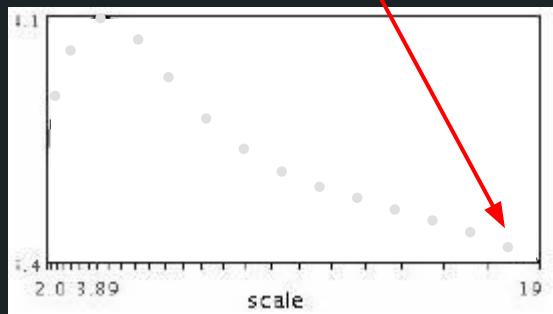
$$f(I_{i \sqcup i_m}(x, \sigma))$$



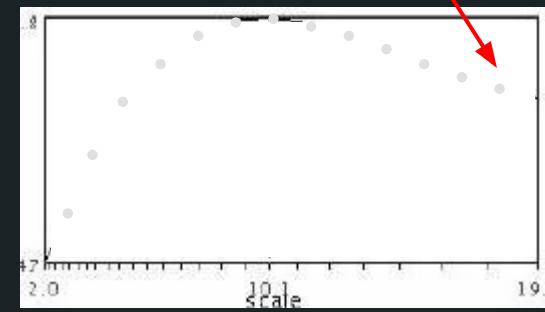
$$f(I_{i \sqcup i_m}(x', \sigma))$$

# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



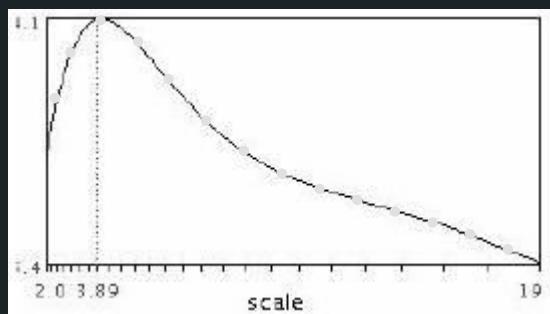
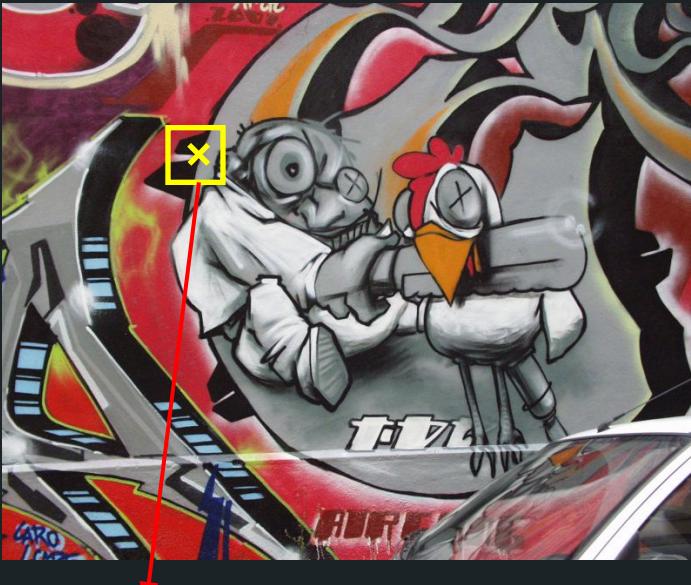
$$f(I_{i \sqcup i_m}(x, \sigma))$$



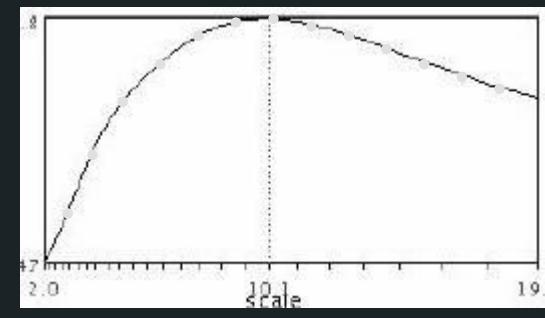
$$f(I_{i \sqcup i_m}(x', \sigma))$$

# Automatic Scale Selection

- Function responses for increasing scale (scale signature)



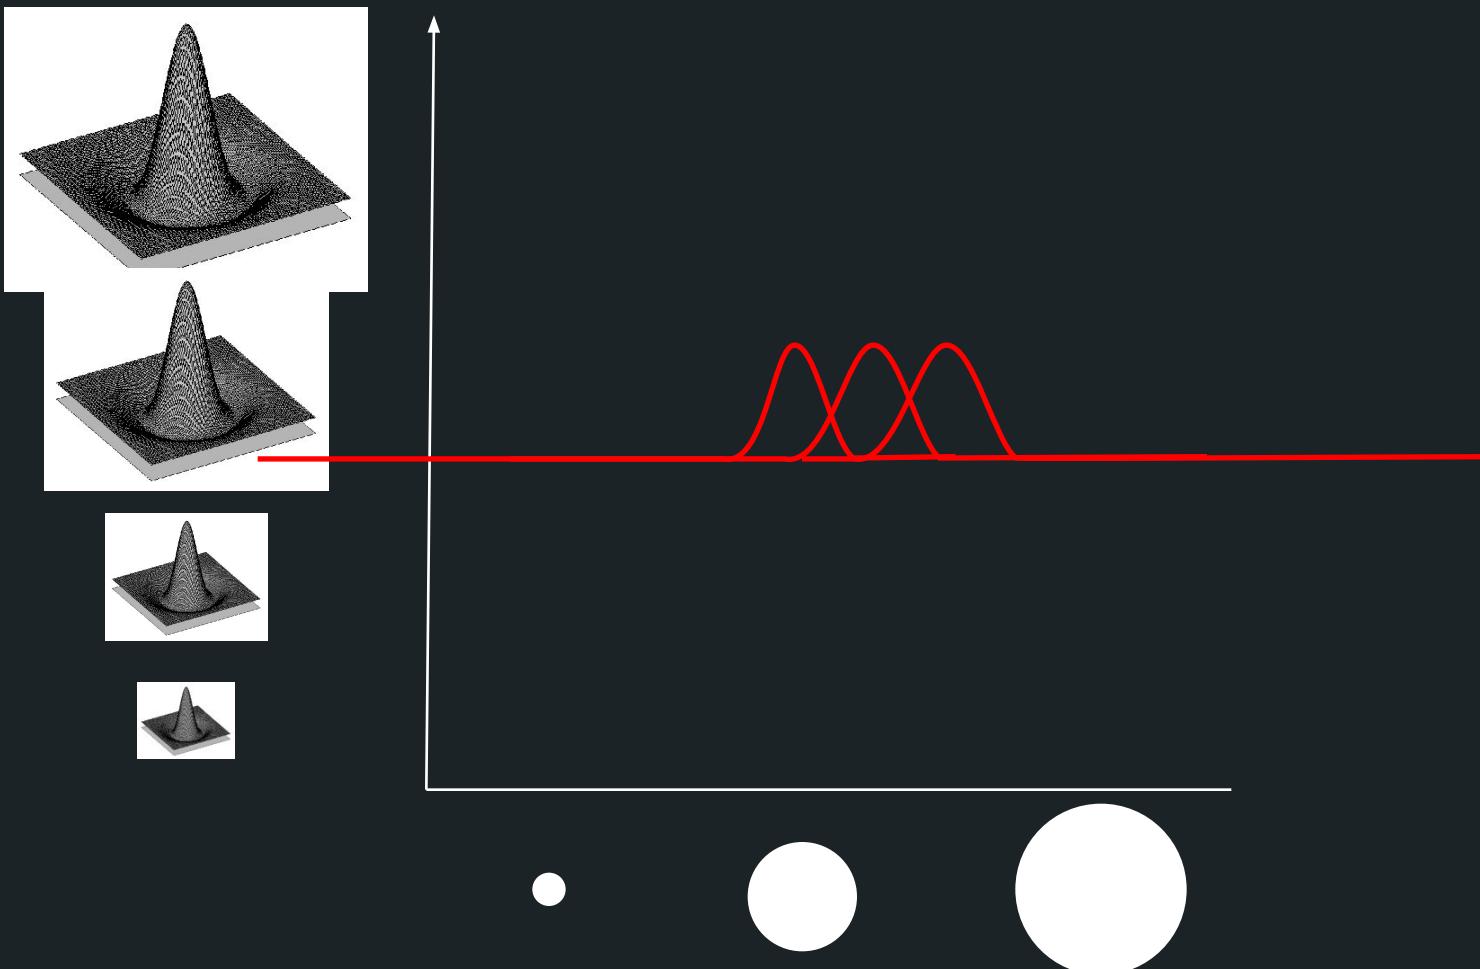
$$f(I_{i \sqcup i_m}(x, \sigma))$$



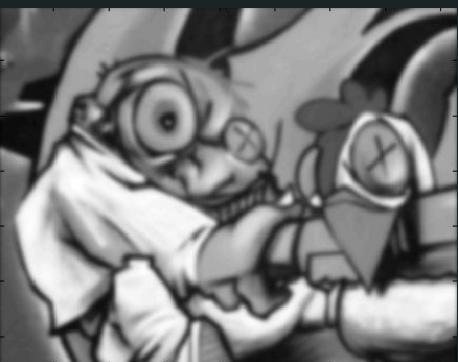
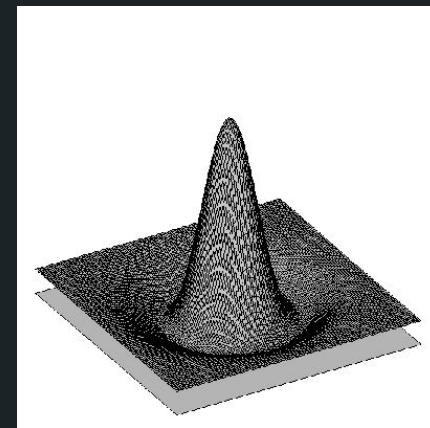
$$f(I_{i \sqcup i_m}(x', \sigma'))$$

# What Is A Useful Signature Function?

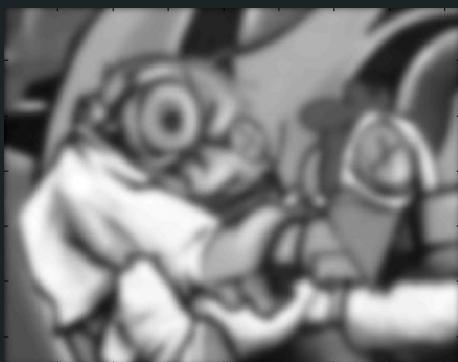
- Difference-of-Gaussian = “blob” detector



# Difference-of-Gaussian (DoG)



-

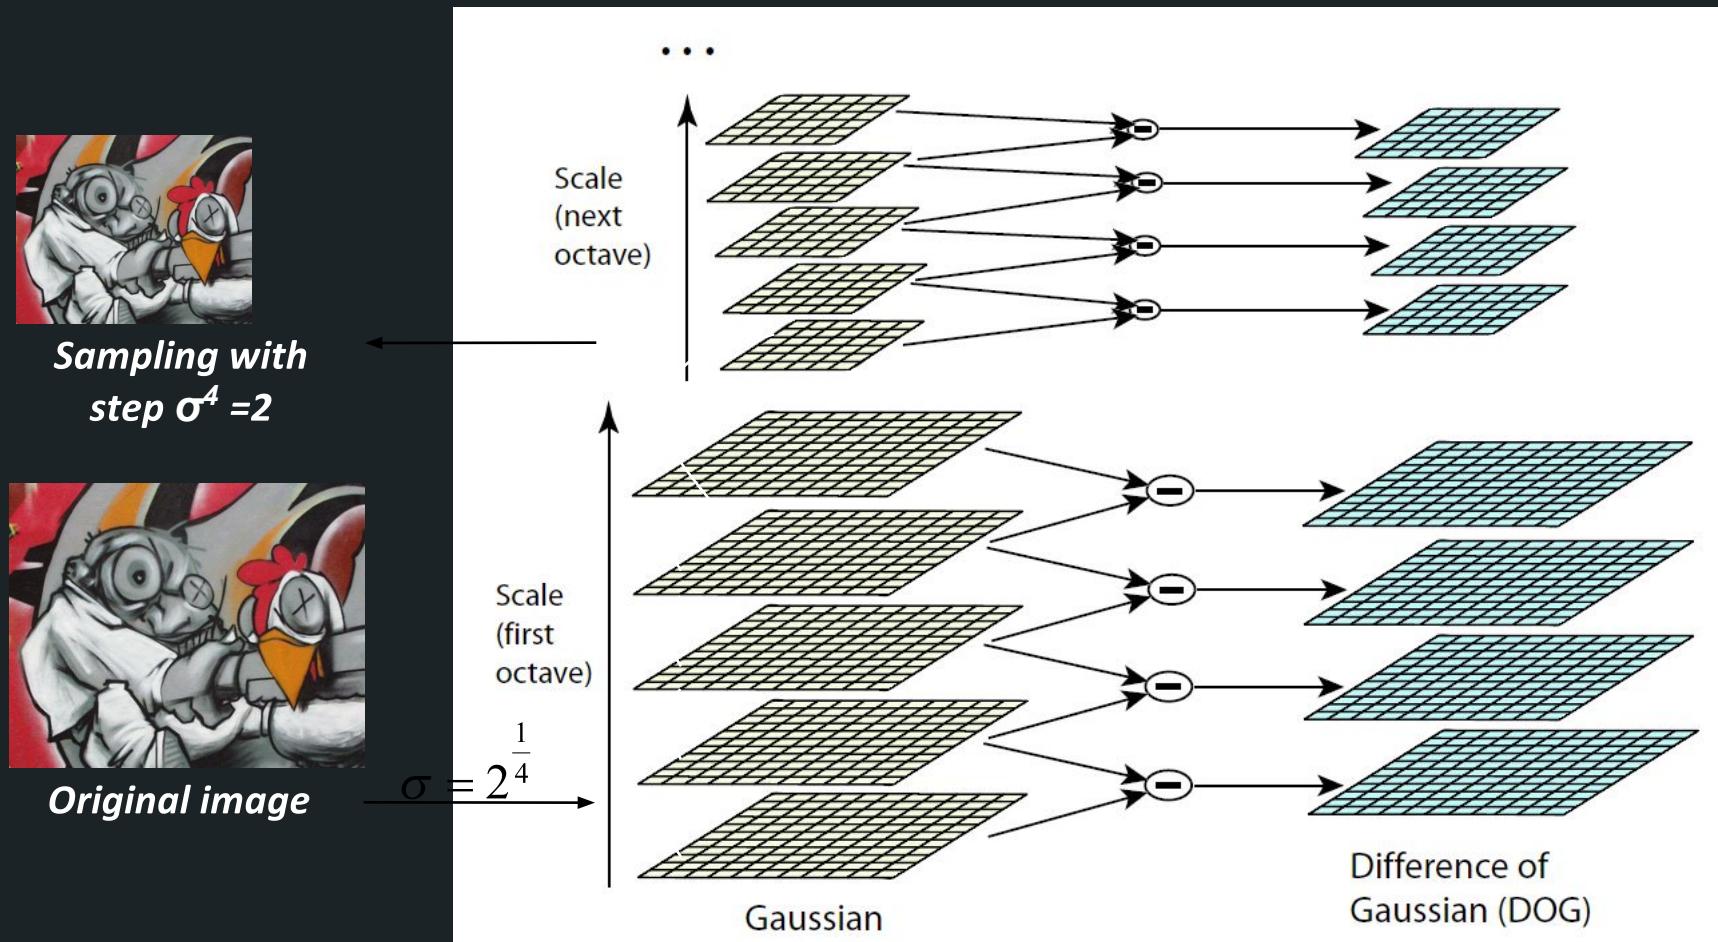


=

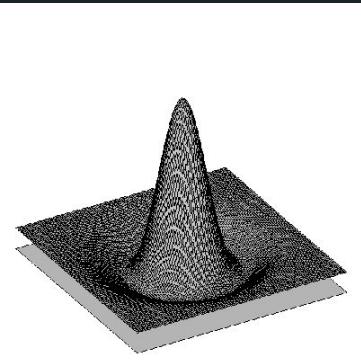


# DoG – Efficient Computation

- Computation in Gaussian scale pyramid

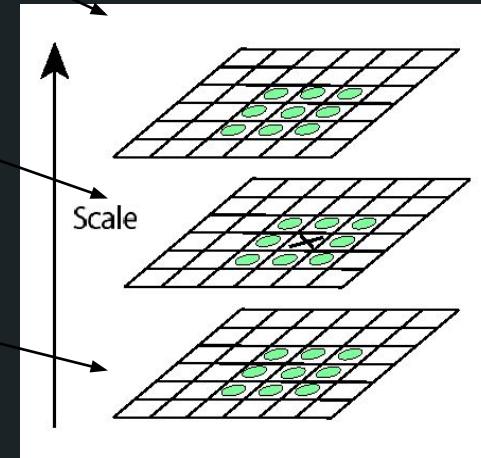
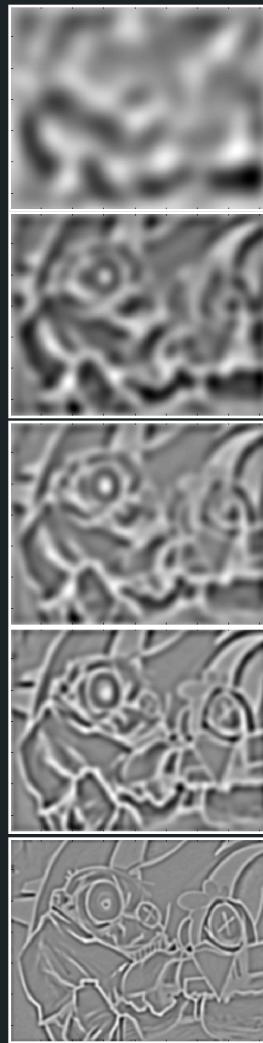


# Find local maxima in position-scale space of Difference-of-Gaussian



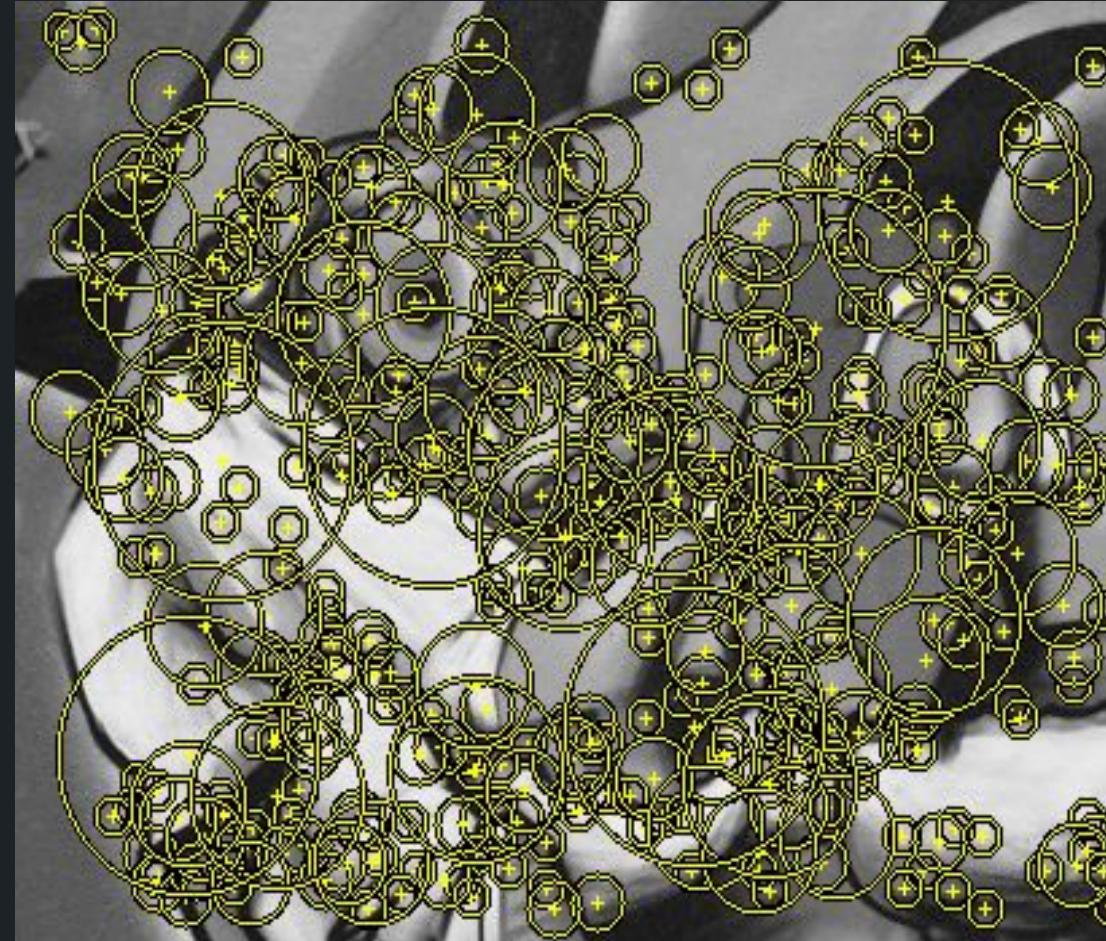
$$L_{xx}(\sigma) + L_{yy}(\sigma) \rightarrow \sigma^3$$

Diagram illustrating the computation of the Difference-of-Gaussian (DoG) operator. A grayscale input image is processed through a series of convolutional kernels, represented by arrows pointing upwards. The first arrow leads to a scale of  $\sigma^3$ . Subsequent arrows lead to scales of  $\sigma^2$ ,  $\sigma^4$ , and  $\sigma^5$ .



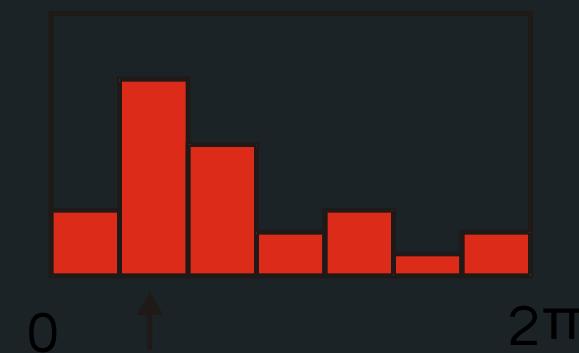
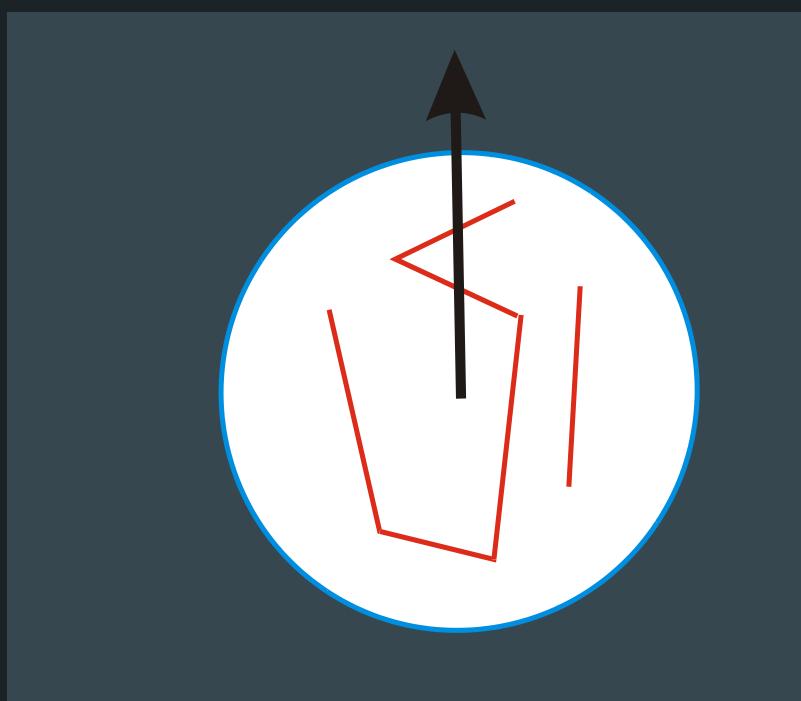
⇒ List of  
 $(x, y, s)$

# Results: Difference-of-Gaussian



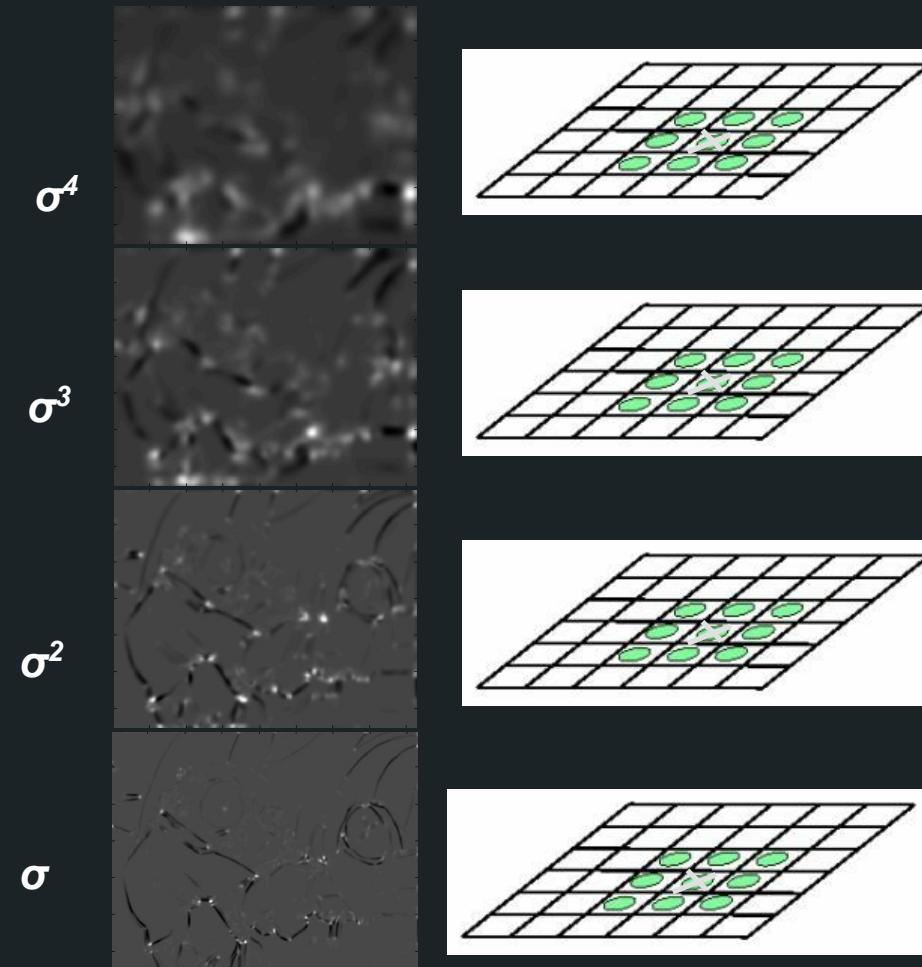
# Orientation Normalization

- Compute orientation histogram
- Select dominant orientation
- Normalize: rotate to fixed orientation



# Harris-Laplace

## 1. Initialization: Multiscale Harris corner detection



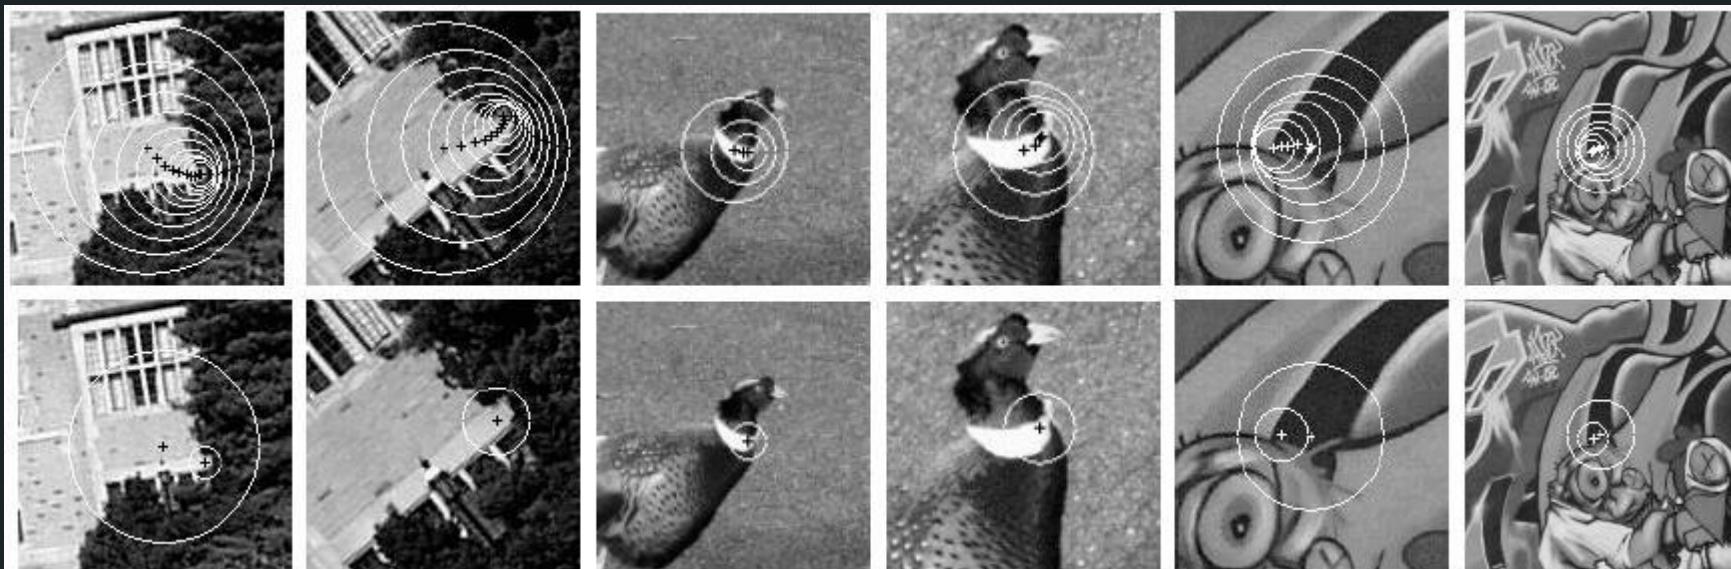
Computing Harris function

Detecting local maxima

# Harris-Laplace

1. Initialization: Multiscale Harris corner detection
2. Scale selection based on Laplacian  
(same procedure with Hessian  $\Rightarrow$  Hessian-Laplace)

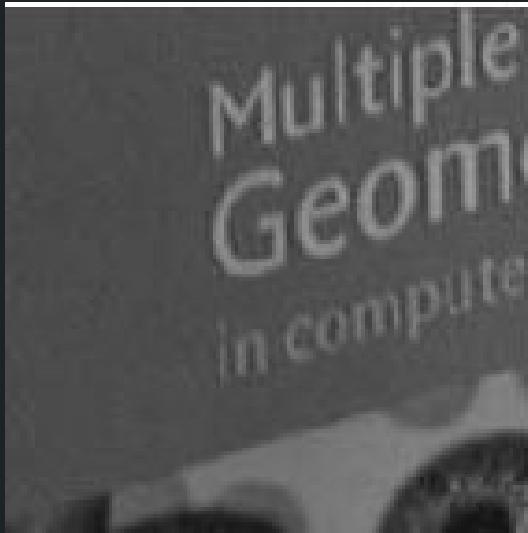
Harris points



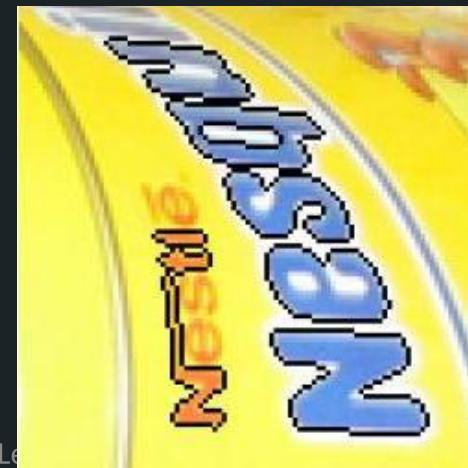
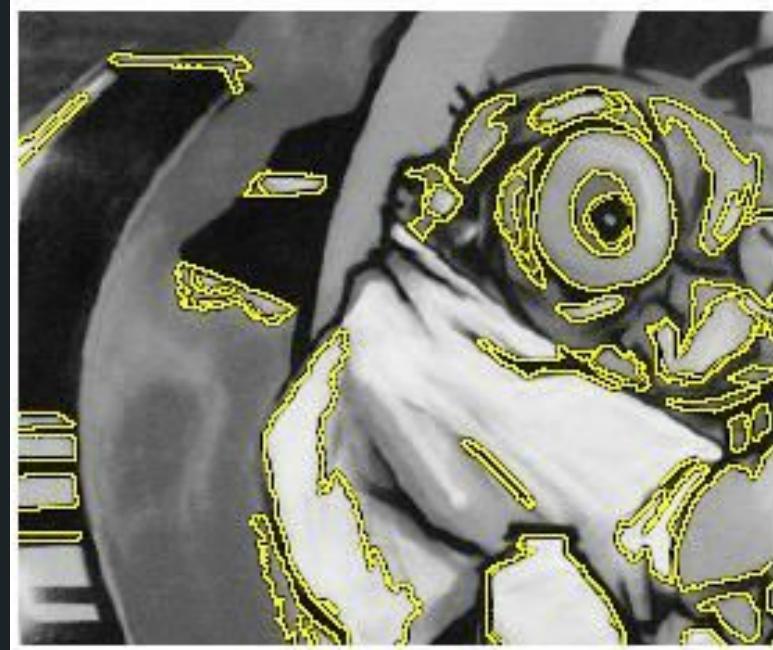
Harris-Laplace points

# Maximally Stable Extremal Regions

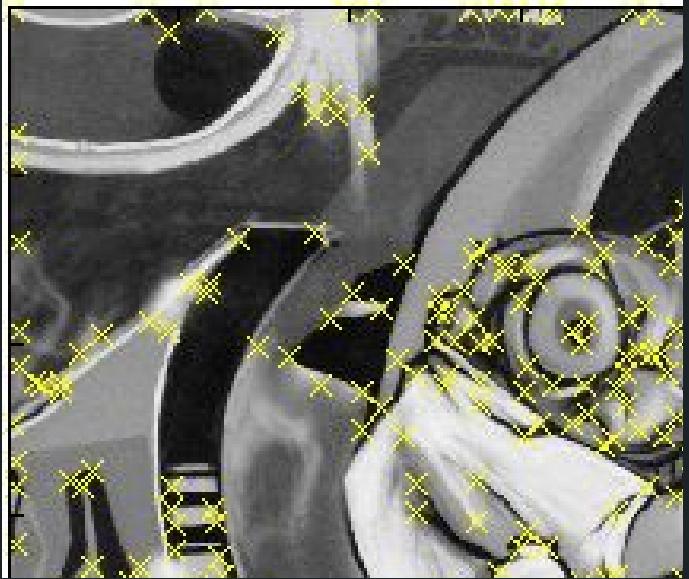
- Based on Watershed segmentation algorithm
- Select regions that stay stable over a large parameter range



# Example Results: MSER



# Comparison



Harris



LoG



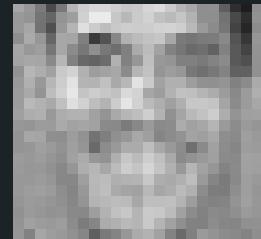
Hessian



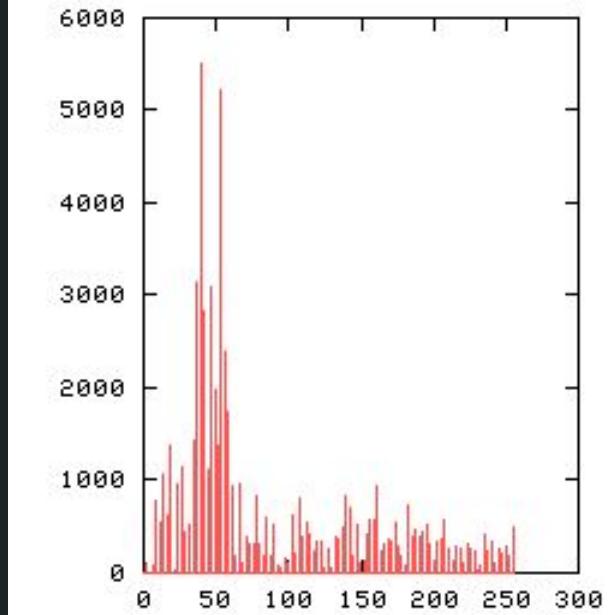
MSER

# Image representations

- Templates
  - Intensity, gradients, etc.
- Histograms
  - Color, texture, SIFT descriptors, etc.



# Image Representations: Histograms



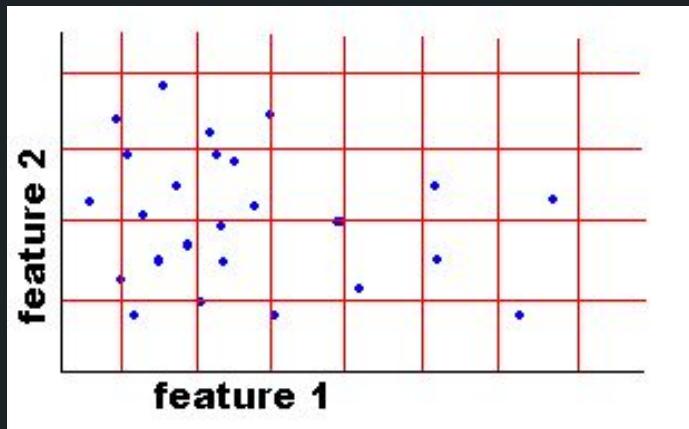
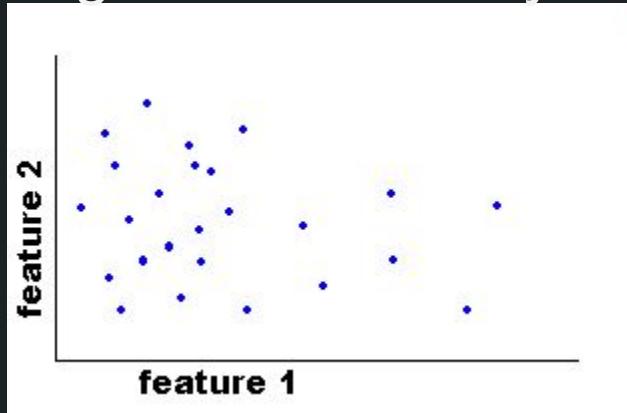
## Global histogram

- Represent distribution of features
  - Color, texture, depth, ...

Space Shuttle  
Cargo Bay

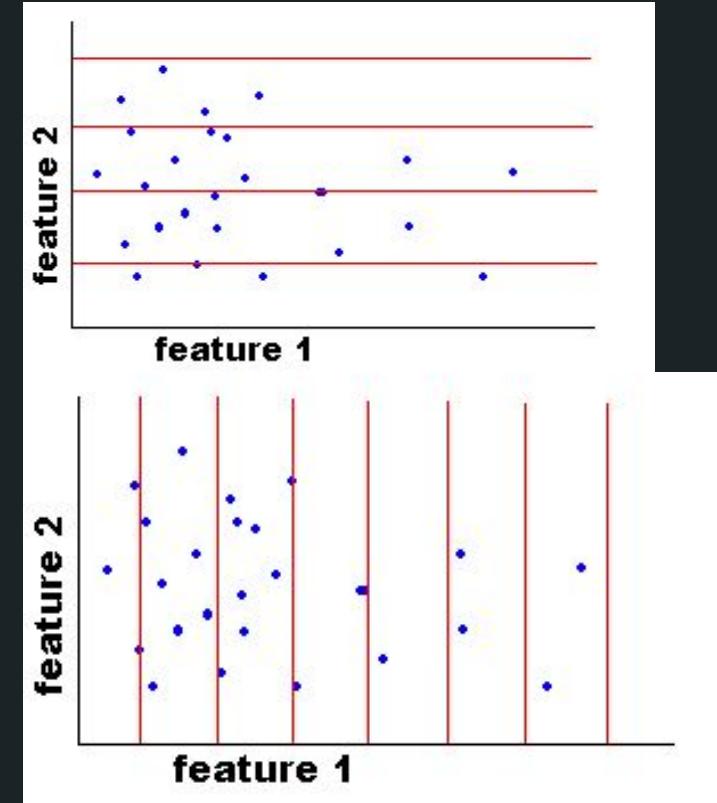
# Image Representations: Histograms

Histogram: Probability or count of data in each bin



- Joint histogram

- Requires lots of data
- Loss of resolution to avoid empty bins

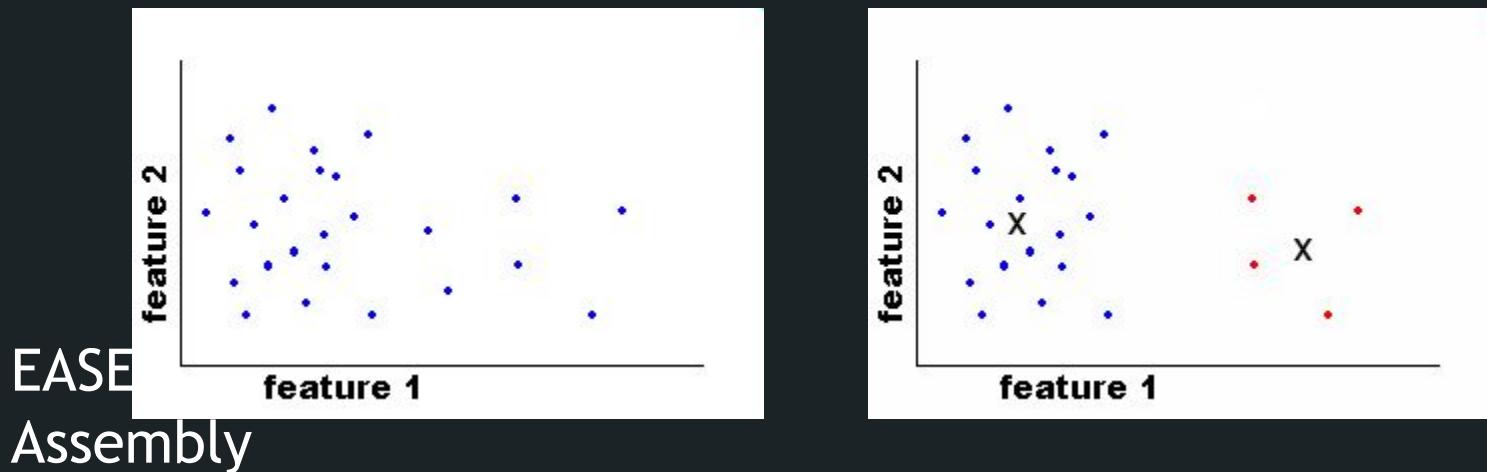


- Marginal histogram

- Requires independent features
- More data/bin than joint histogram

# Image Representations: Histograms

## Clustering



Use the same cluster centers for all images

Space Shuttle  
Cargo Bay

# Computing histogram distance

$$\text{histint}(h_i, h_j) = 1 - \sum_{m=1}^K \min(h_i(m), h_j(m))$$

Histogram intersection (assuming normalized histograms)

$$\chi^2(h_i, h_j) = \frac{1}{2} \sum_{m=1}^K \frac{[h_i(m) - h_j(m)]^2}{h_i(m) + h_j(m)}$$

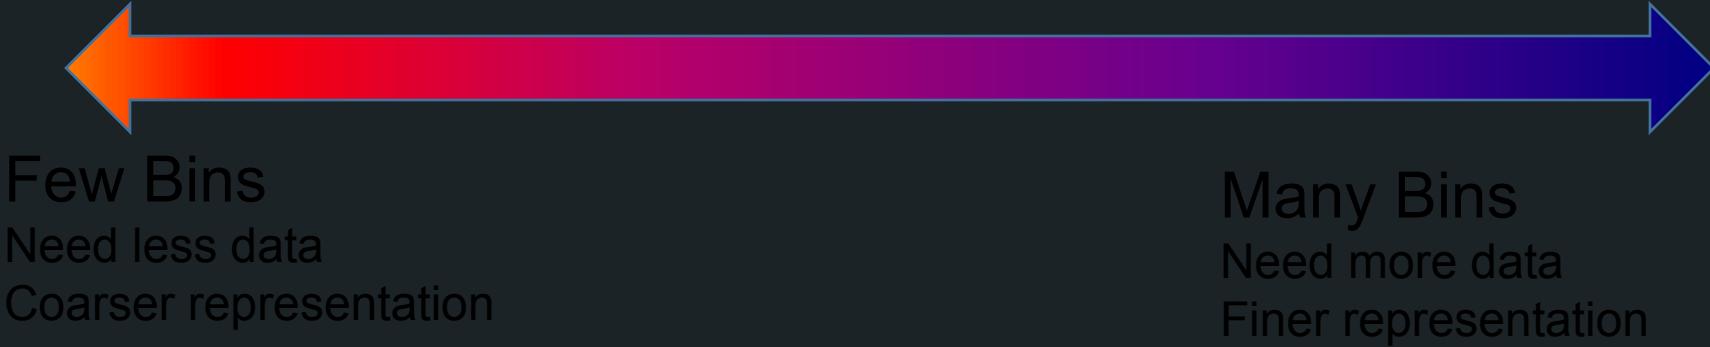
Chi-squared Histogram matching distance



Cars found by color histogram matching using chi-squared

# Histograms: Implementation issues

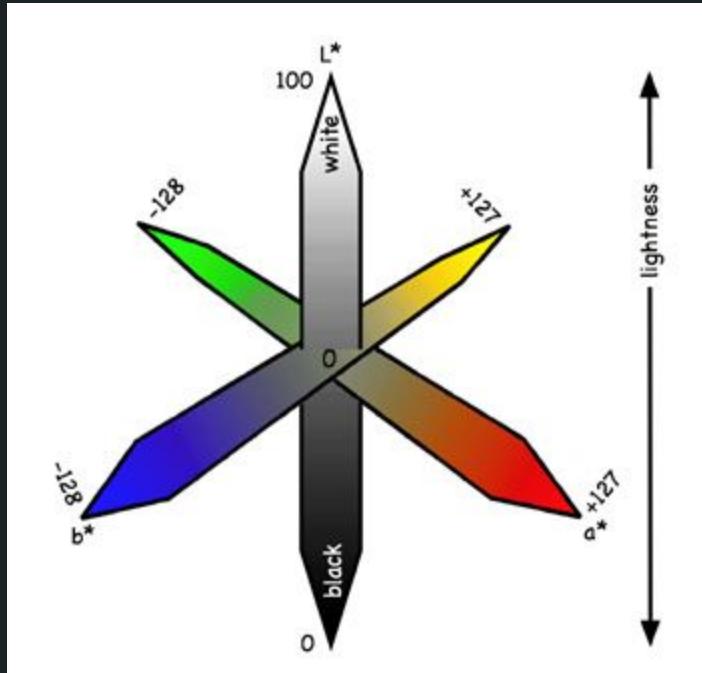
- Quantization
  - Grids: fast but applicable only with few dimensions
  - Clustering: slower but can quantize data in higher dimensions



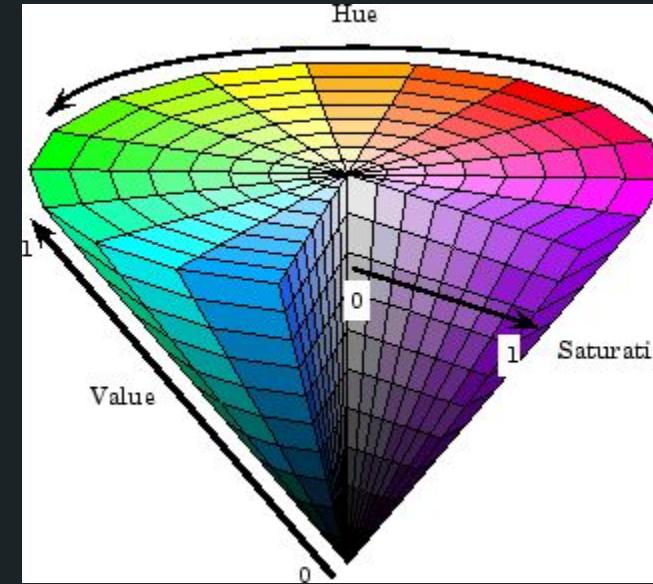
- Matching
  - Histogram intersection or Euclidean may be faster
  - Chi-squared often works better
  - Earth mover's distance is good for when nearby bins represent similar values

# What kind of things do we compute histograms of?

- Color



L<sup>\*</sup>a<sup>\*</sup>b<sup>\*</sup> color space

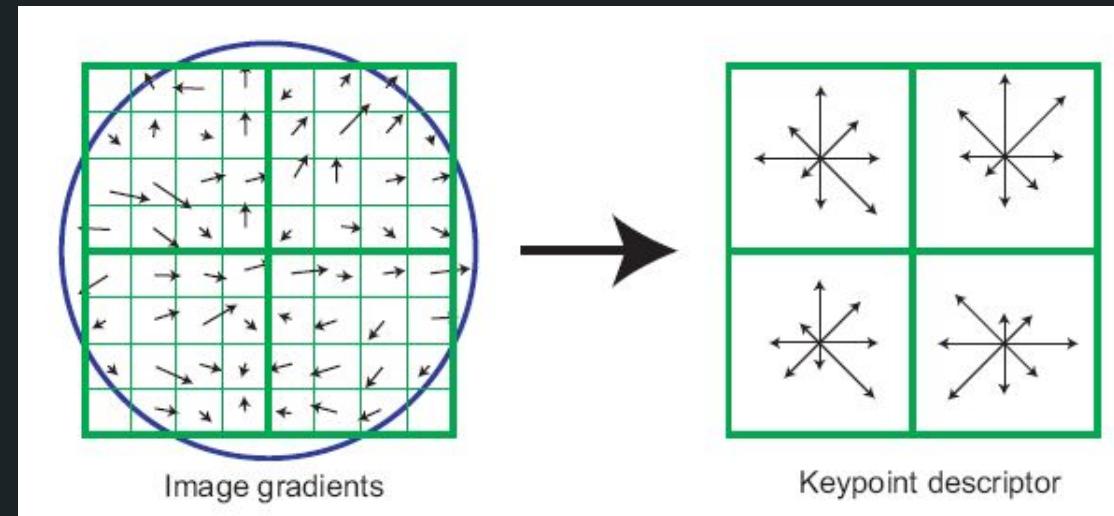


HSV color space

- Texture (filter banks or HOG over regions)

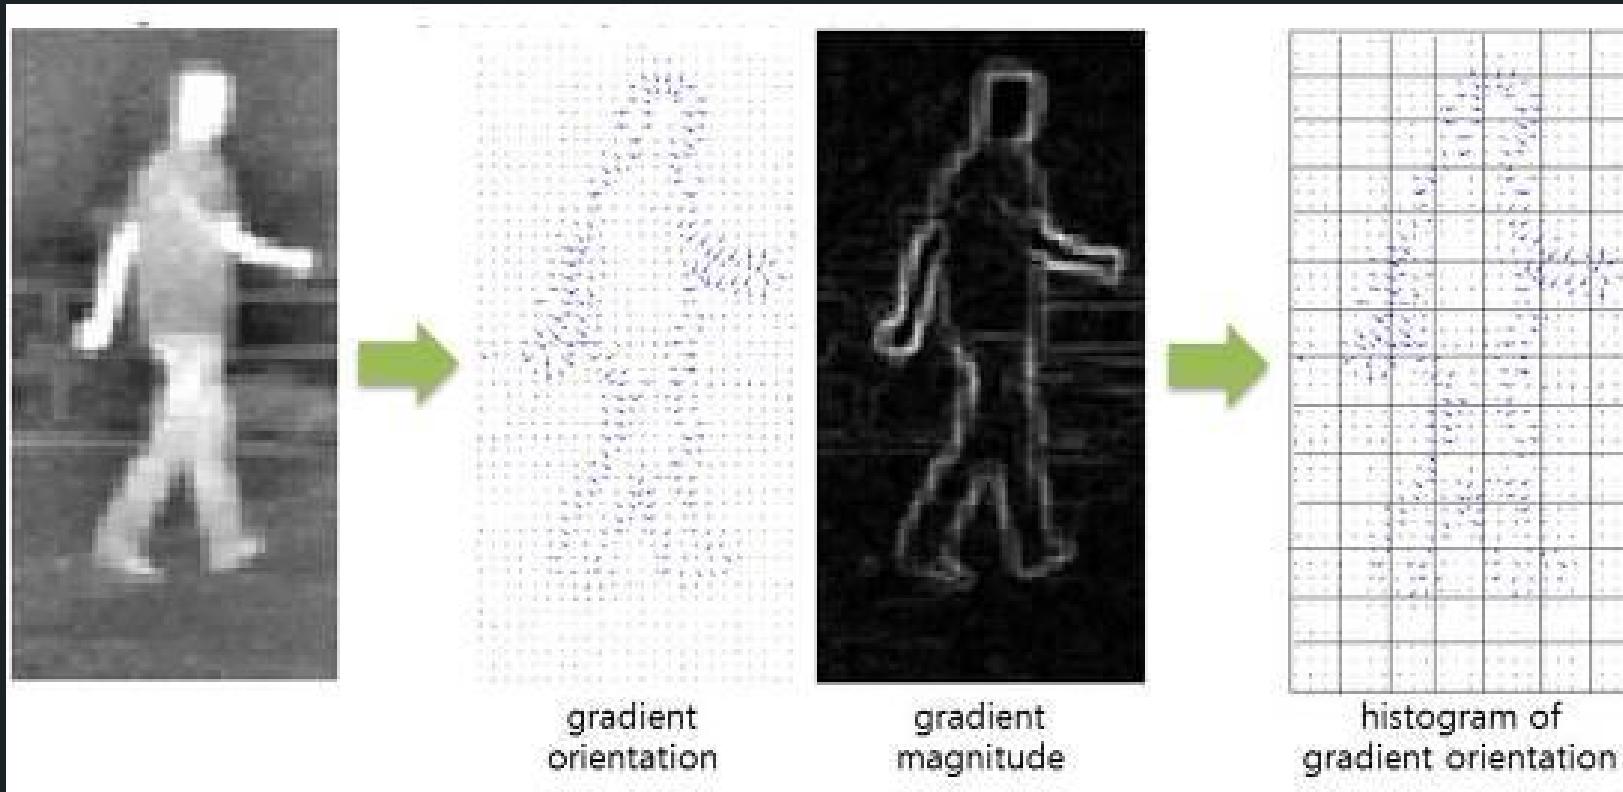
# What kind of things do we compute histograms of?

- Histograms of oriented gradients



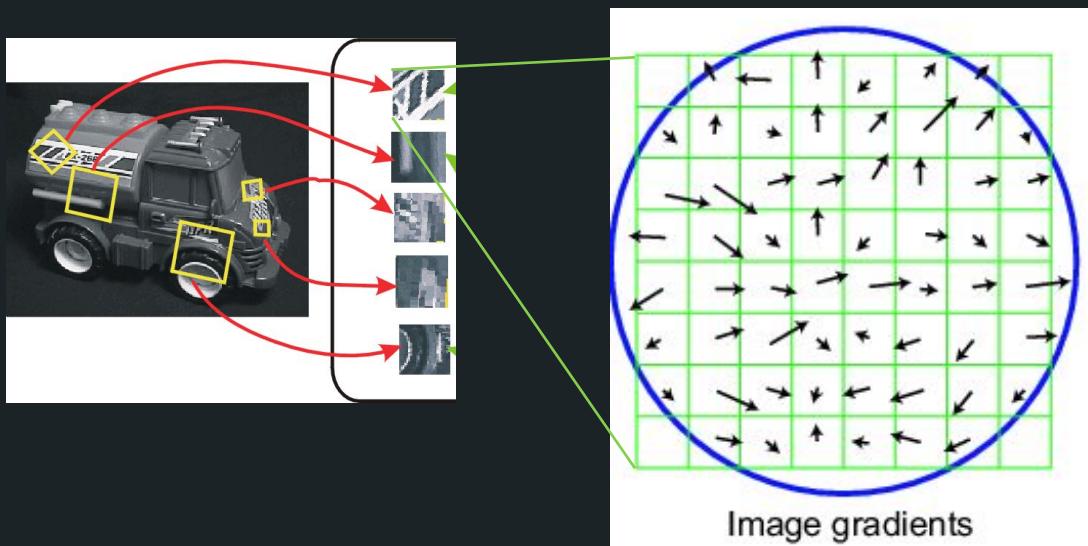
# HoG(Histogram of Oriented Gradients)

Uses gradients(indirectly the edges and so the shape of objects)



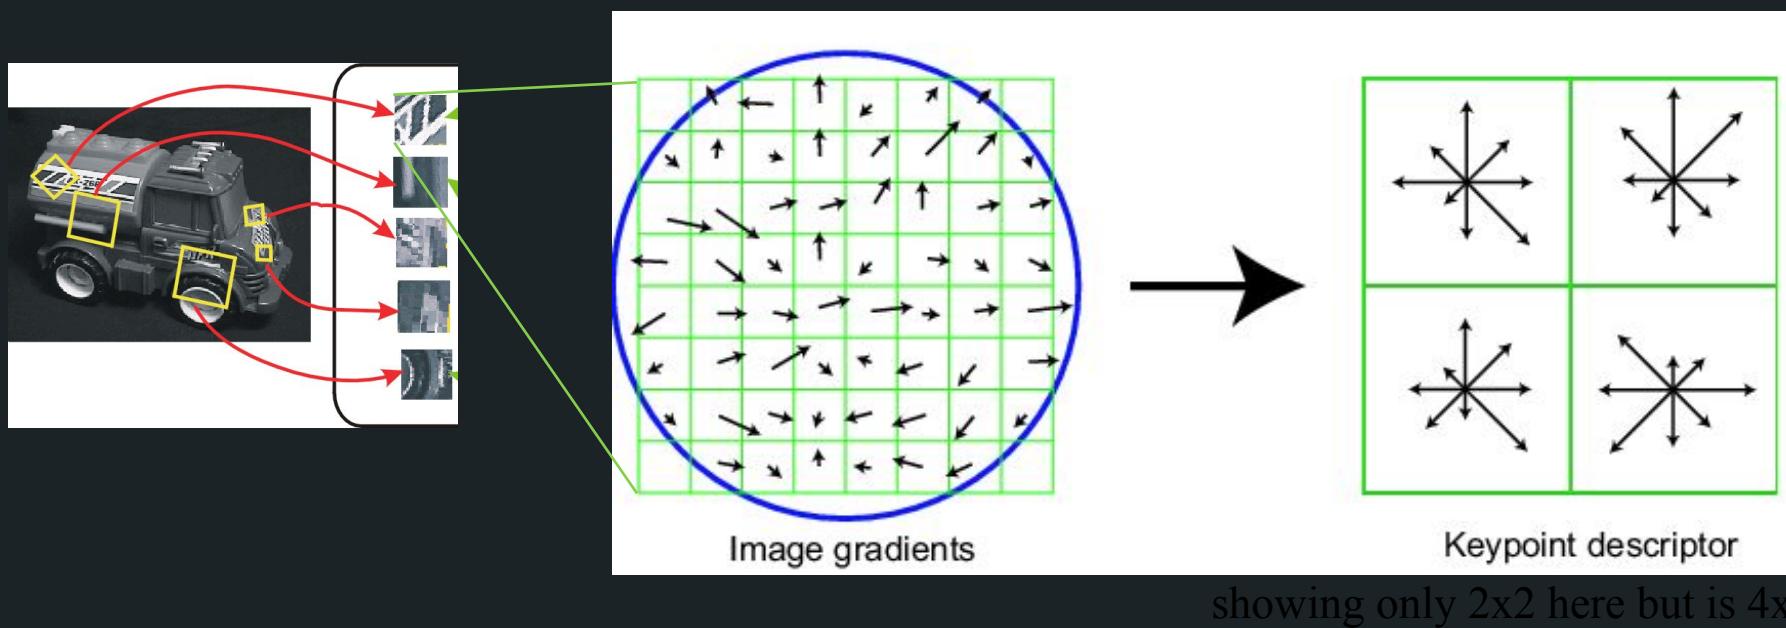
# SIFT vector formation

- Computed on rotated and scaled version of window according to computed orientation & scale
  - resample the window
- Based on gradients weighted by a Gaussian of variance half the window (for smooth falloff)



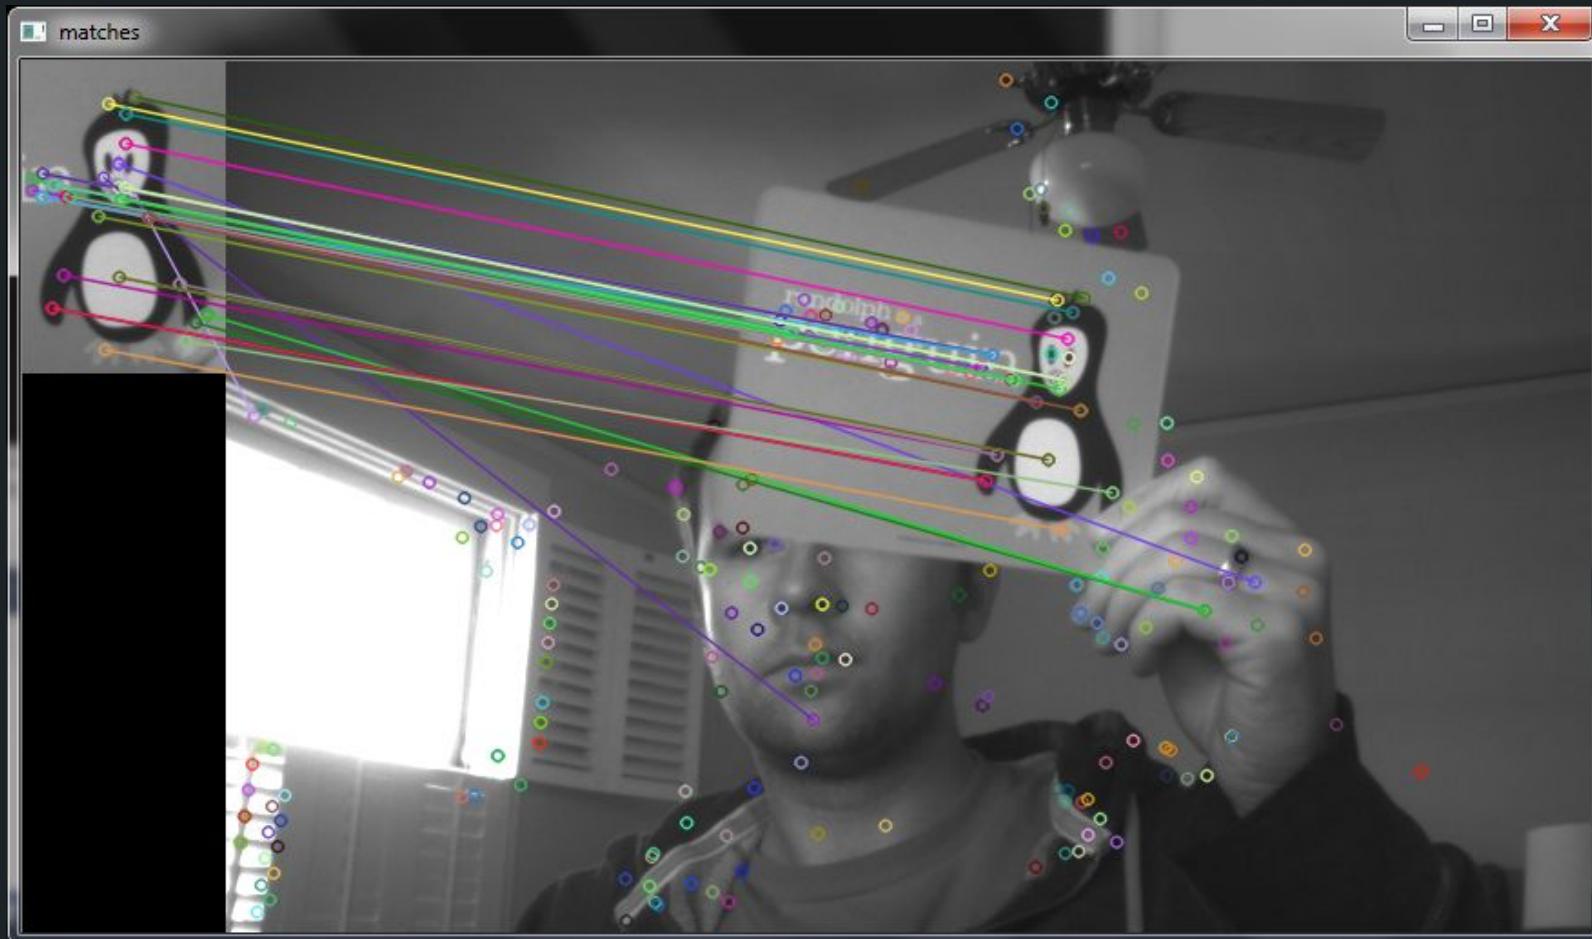
# SIFT vector formation

- 4x4 array of gradient orientation histogram weighted by magnitude
- 8 orientations x 4x4 array = 128 dimensions
- Motivation: some sensitivity to spatial layout, but not too much.



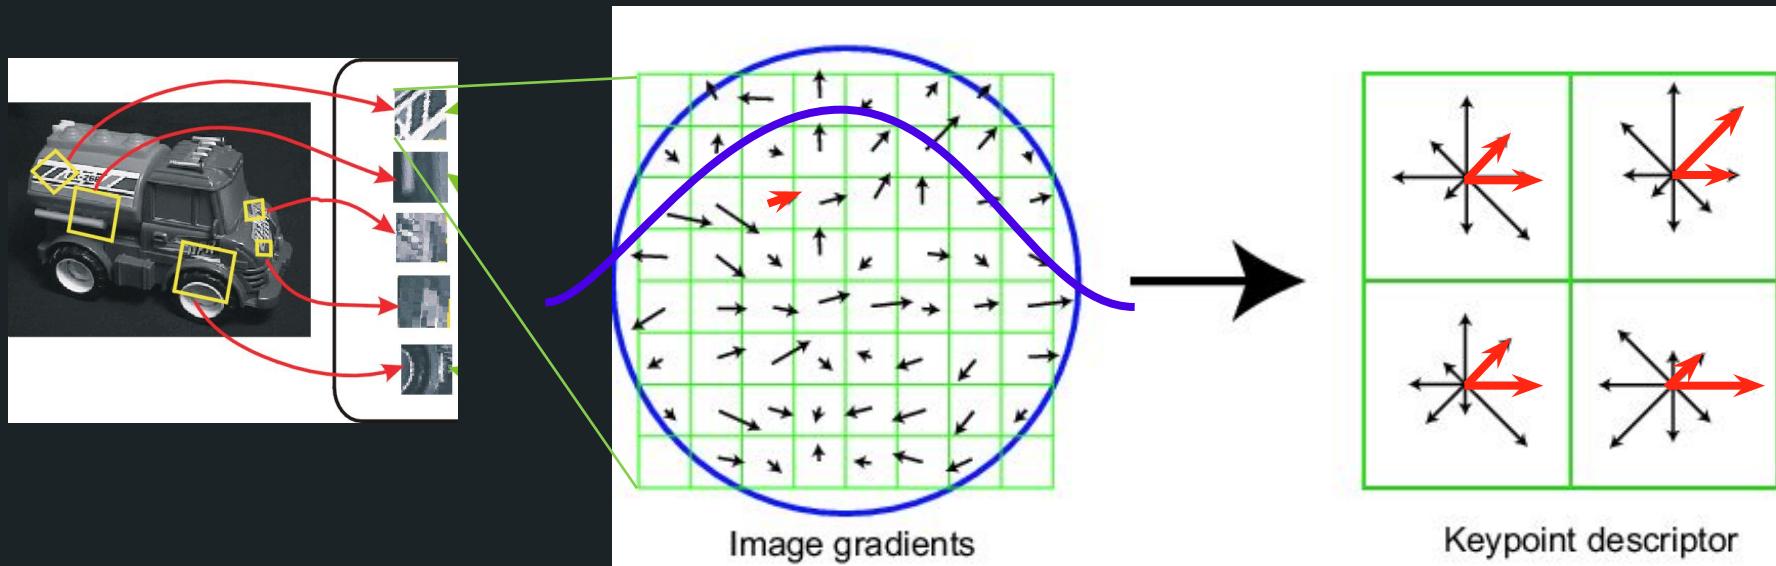
# SIFT (Scale-invariant feature transform)

Invariant to image translation , scaling and rotation



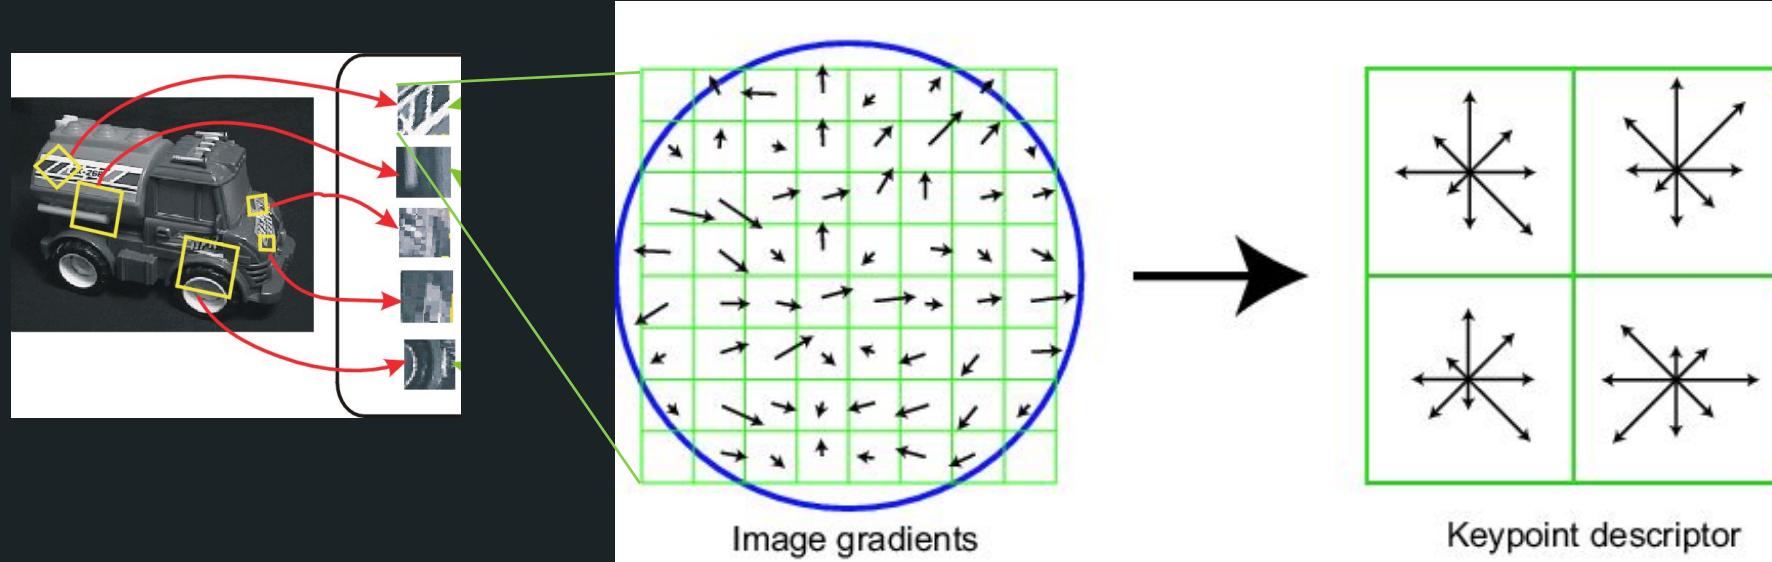
# Ensure smoothness

- Gaussian weight
- Trilinear interpolation
  - a given gradient contributes to 8 bins:  
4 in space times 2 in orientation

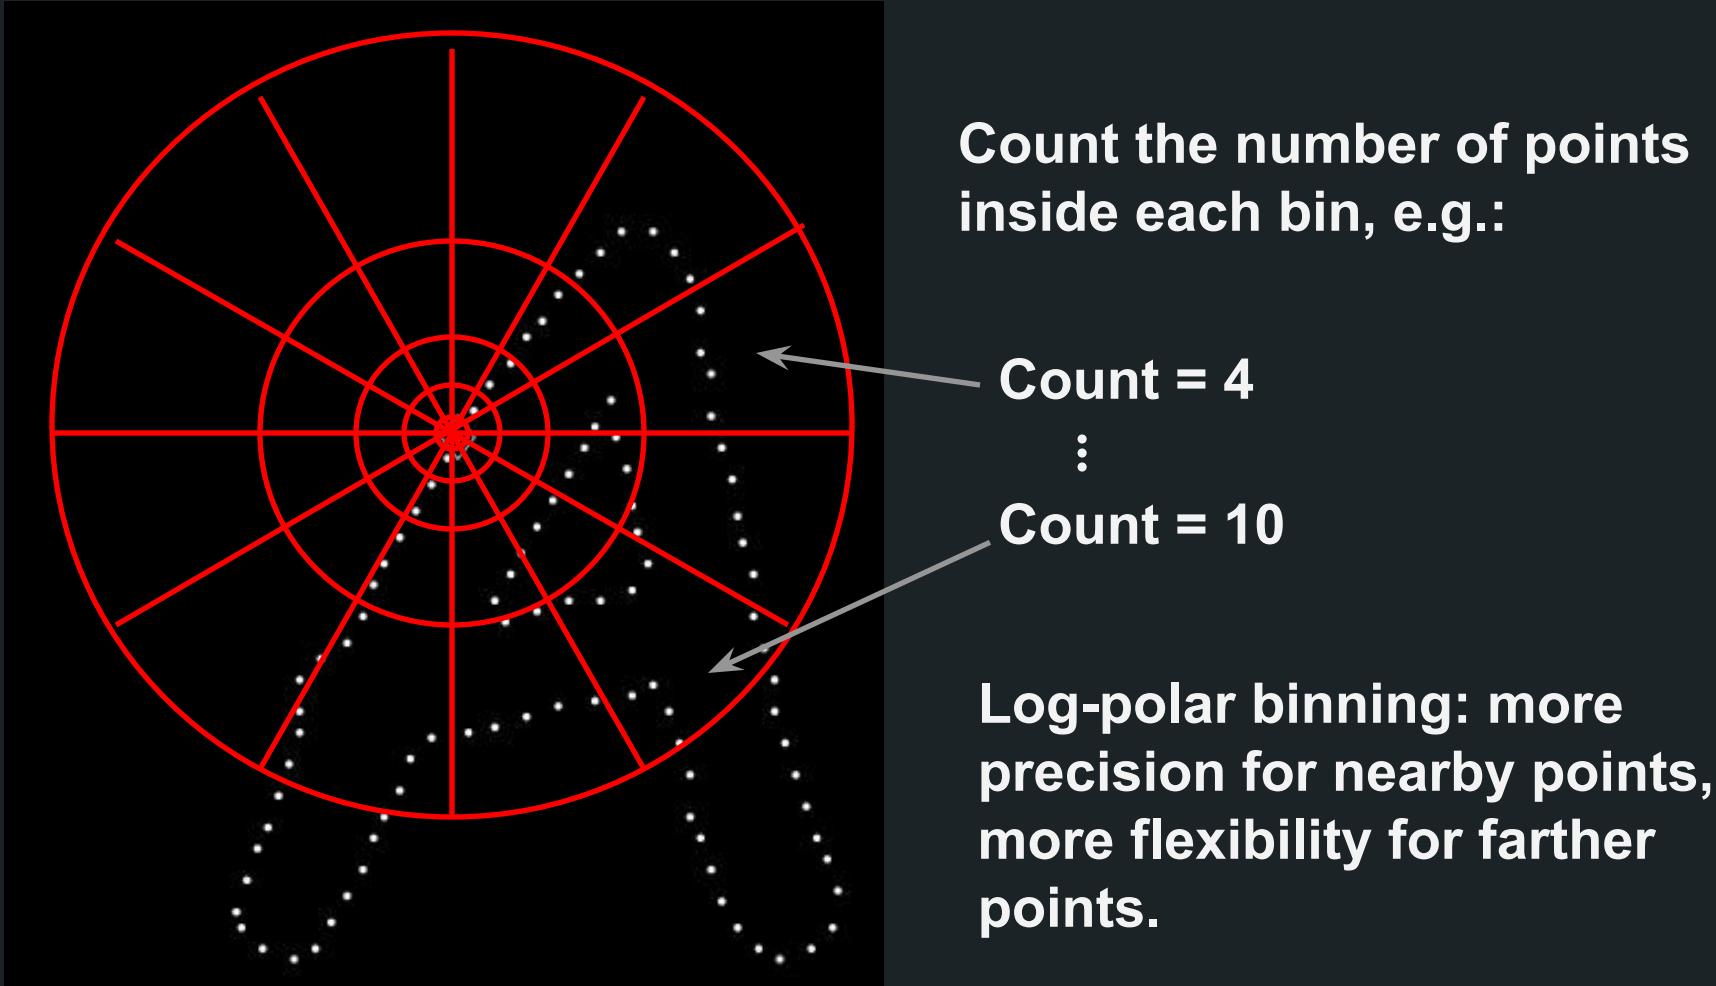


# Reduce effect of illumination

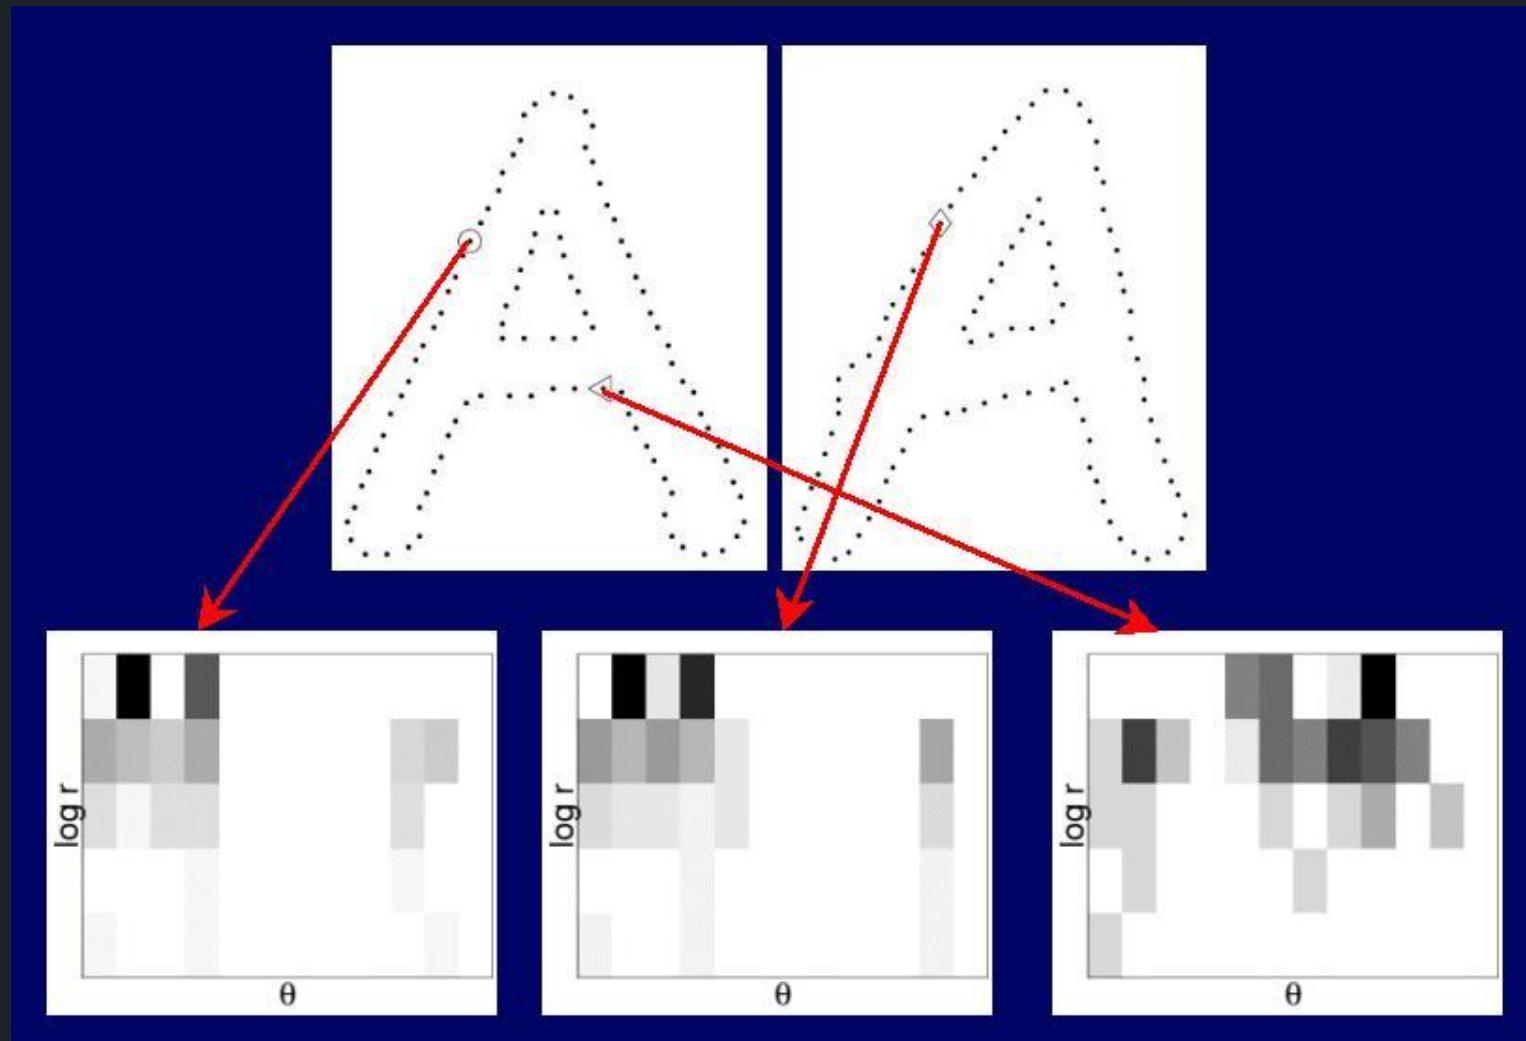
- 128-dim vector normalized to 1
- Threshold gradient magnitudes to avoid excessive influence of high gradients
  - after normalization, clamp gradients  $>0.2$
  - renormalize



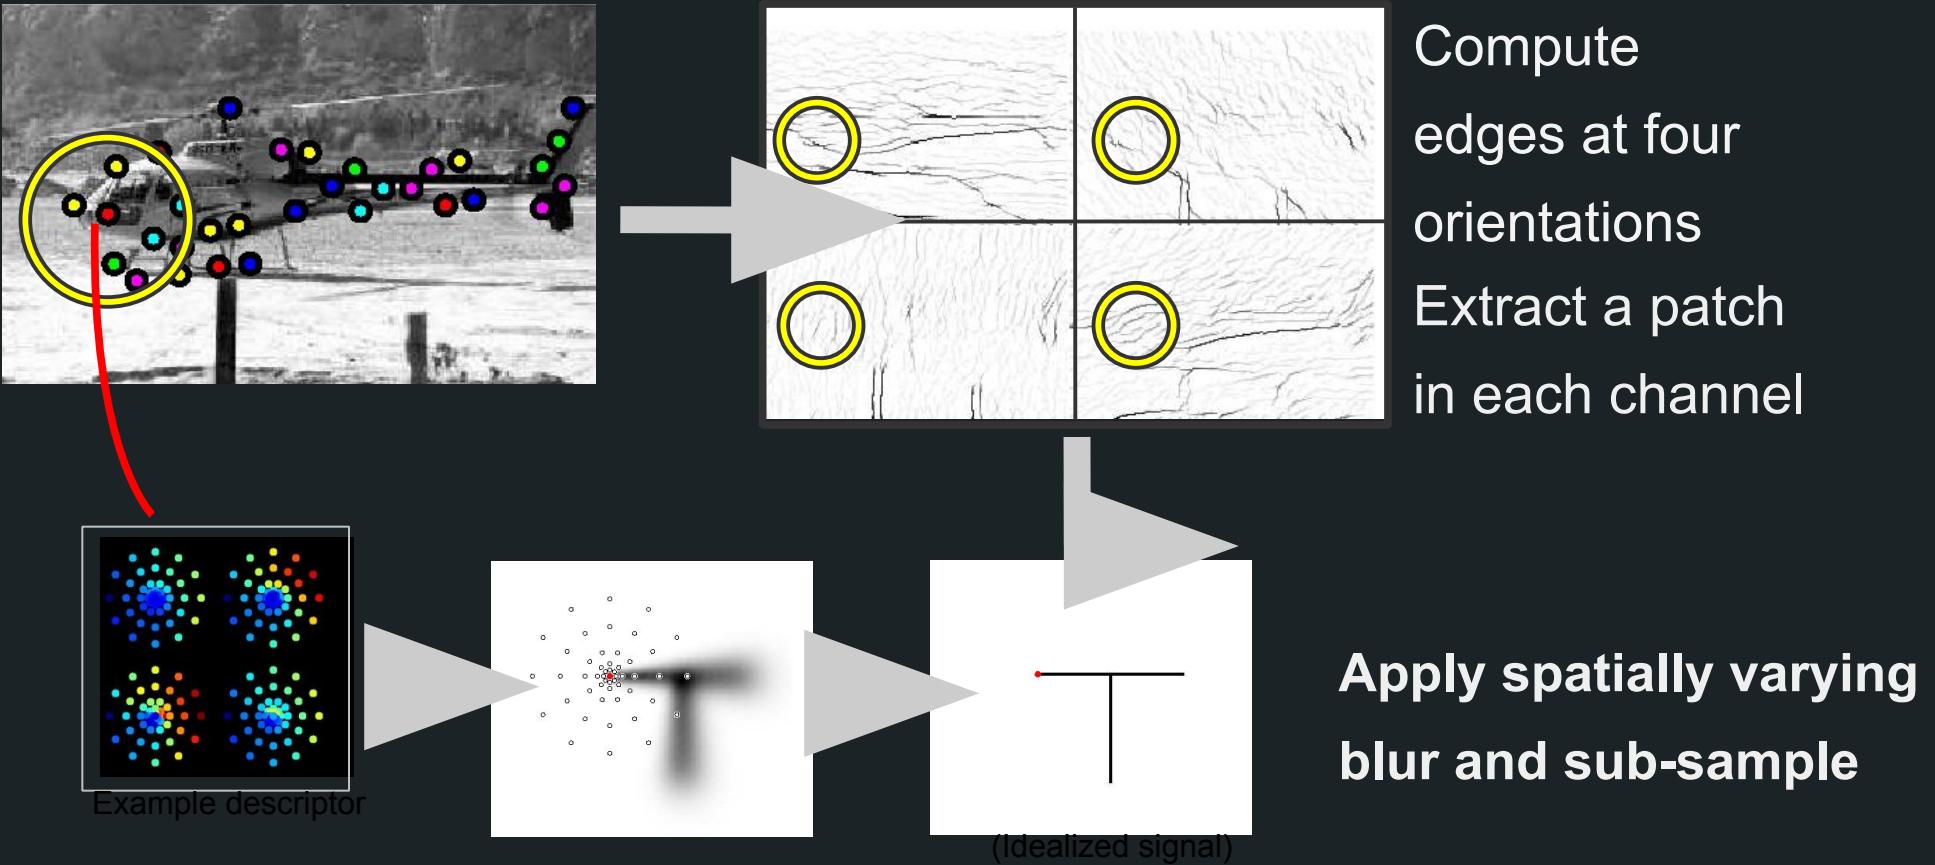
# Local Descriptors: Shape Context



# Shape Context Descriptor



# Local Descriptors: Geometric Blur

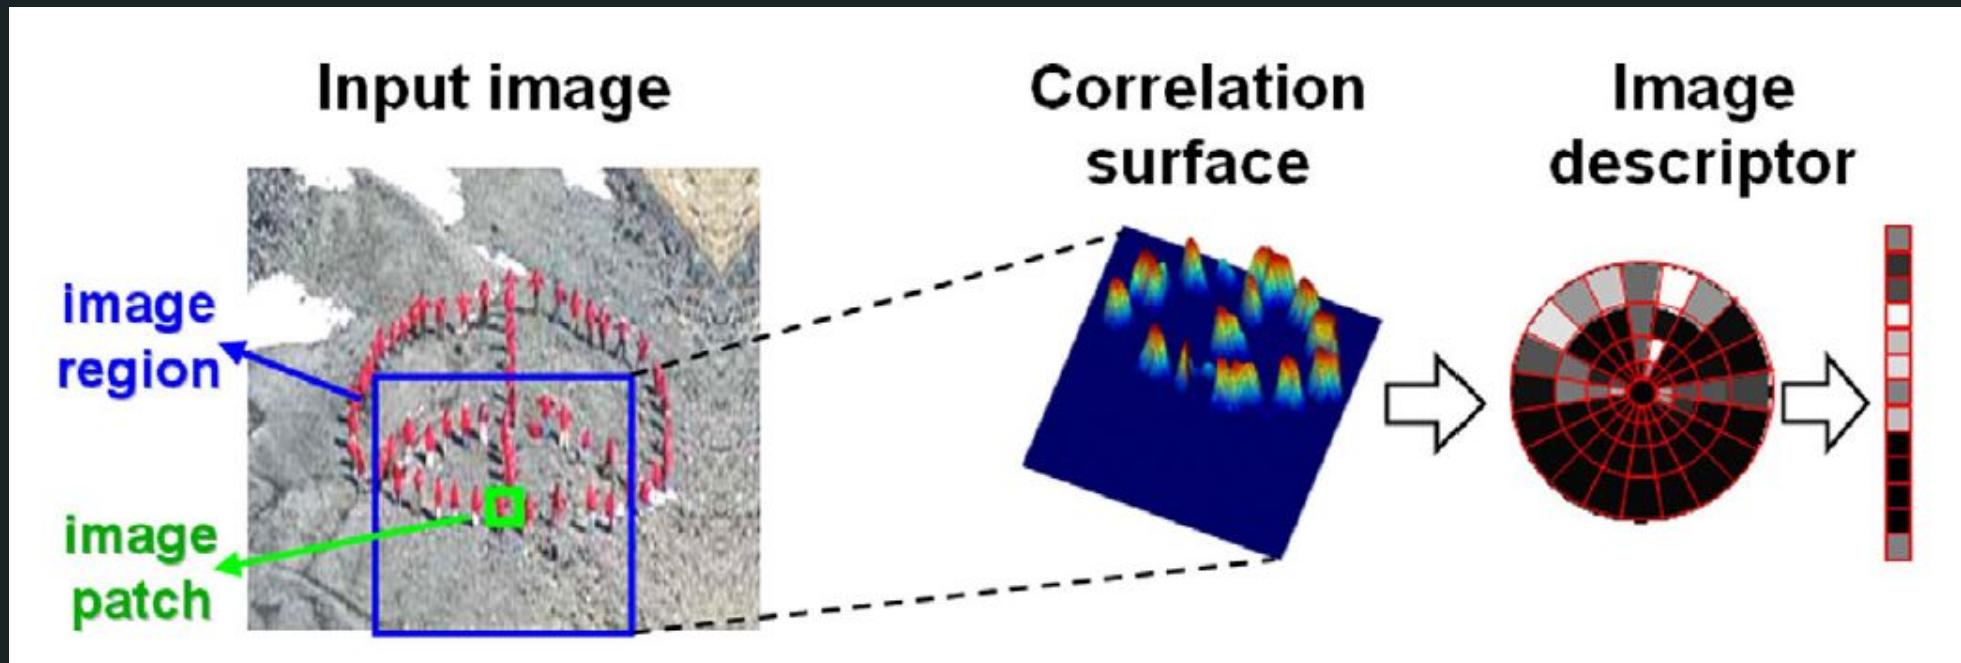


# Self-similarity Descriptor

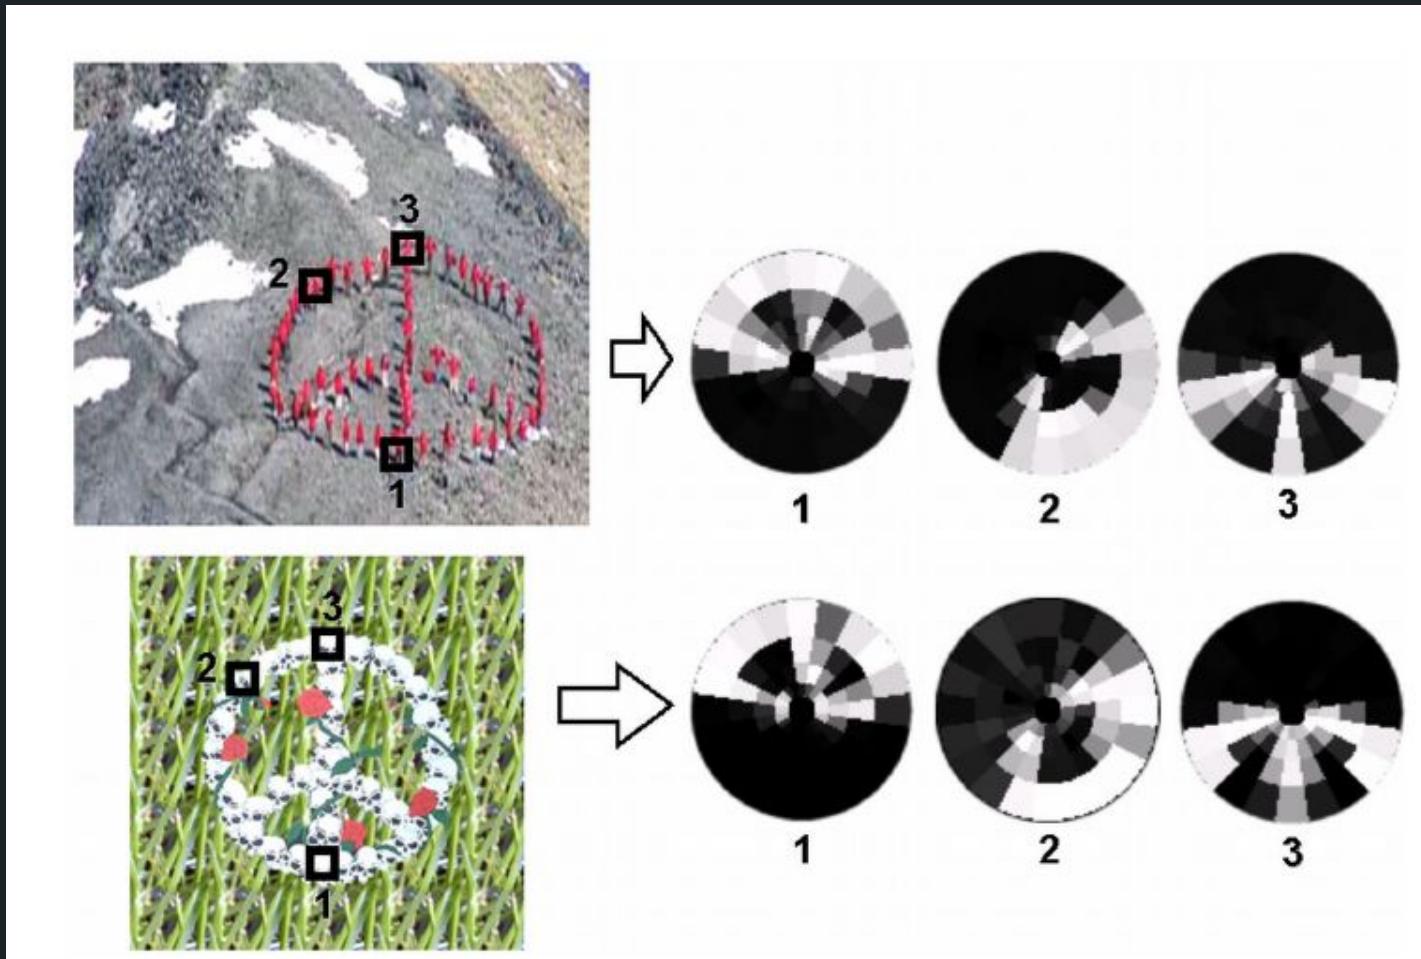


Figure 1. *These images of the same object (a heart) do NOT share common image properties (colors, textures, edges), but DO share a similar geometric layout of local internal self-similarities.*

# Self-similarity Descriptor



# Self-similarity Descriptor



# Right features depend on what you want to know

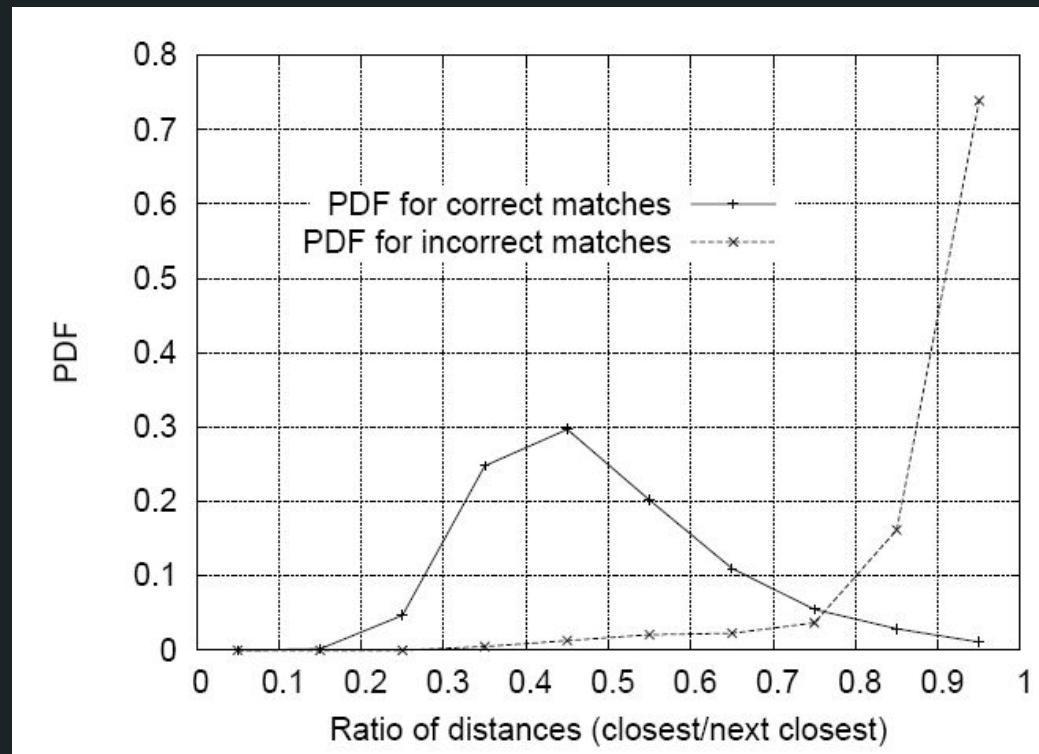
- Shape: scene-scale, object-scale, detail-scale
  - 2D form, shading, shadows, texture, linear perspective
- Material properties: albedo, feel, hardness, ...
  - Color, texture
- Motion
  - Optical flow, tracked points
- Distance
  - Stereo, position, occlusion, scene shape
  - If known object: size, other objects

# Local Descriptors

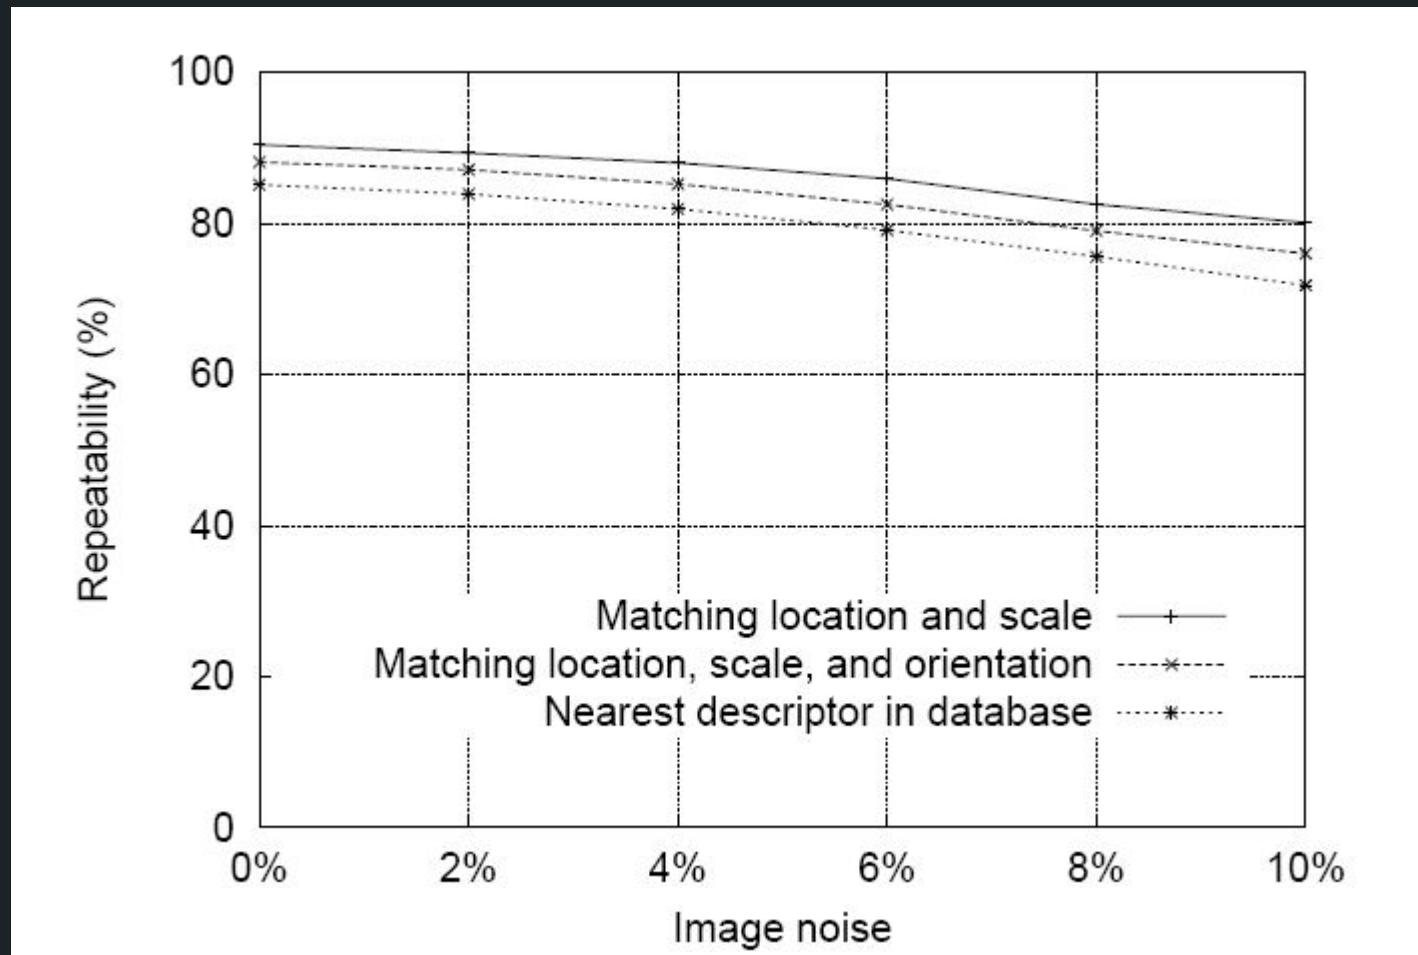
- Most features can be thought of as templates, histograms (counts), or combinations
- The ideal descriptor should be
  - Robust
  - Distinctive
  - Compact
  - Efficient
- Most available descriptors focus on edge/gradient information
  - Capture texture information
  - Color rarely used

# Matching Local Features

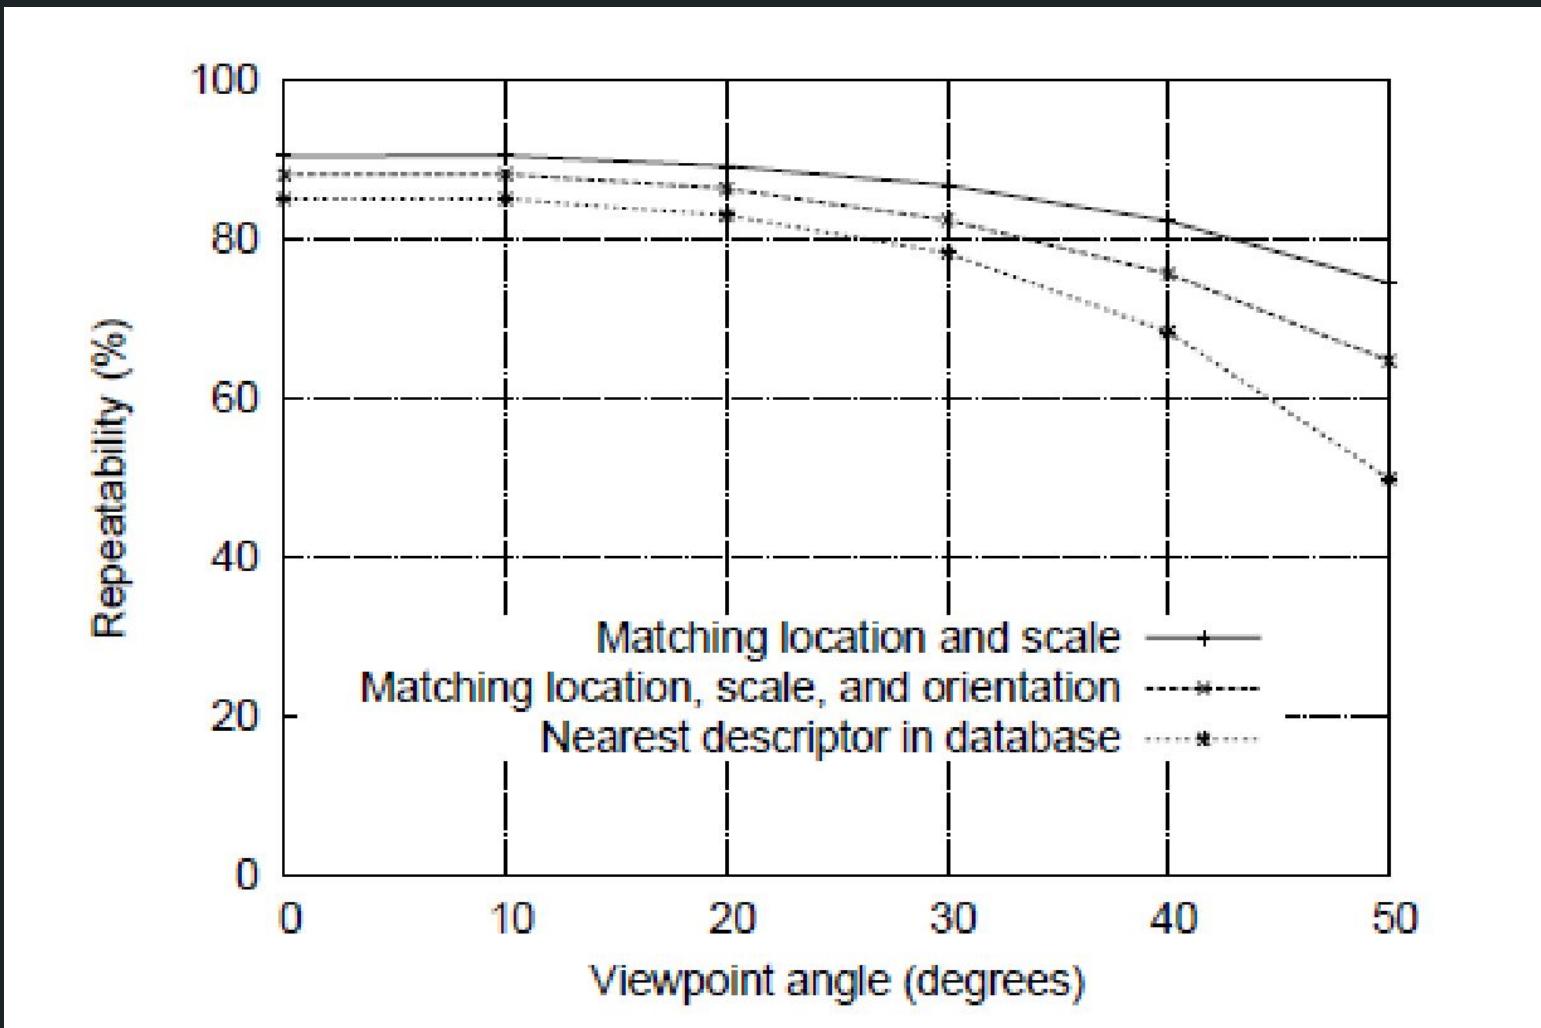
- Nearest neighbor (Euclidean distance)
- Threshold ratio of nearest to 2<sup>nd</sup> nearest descriptor



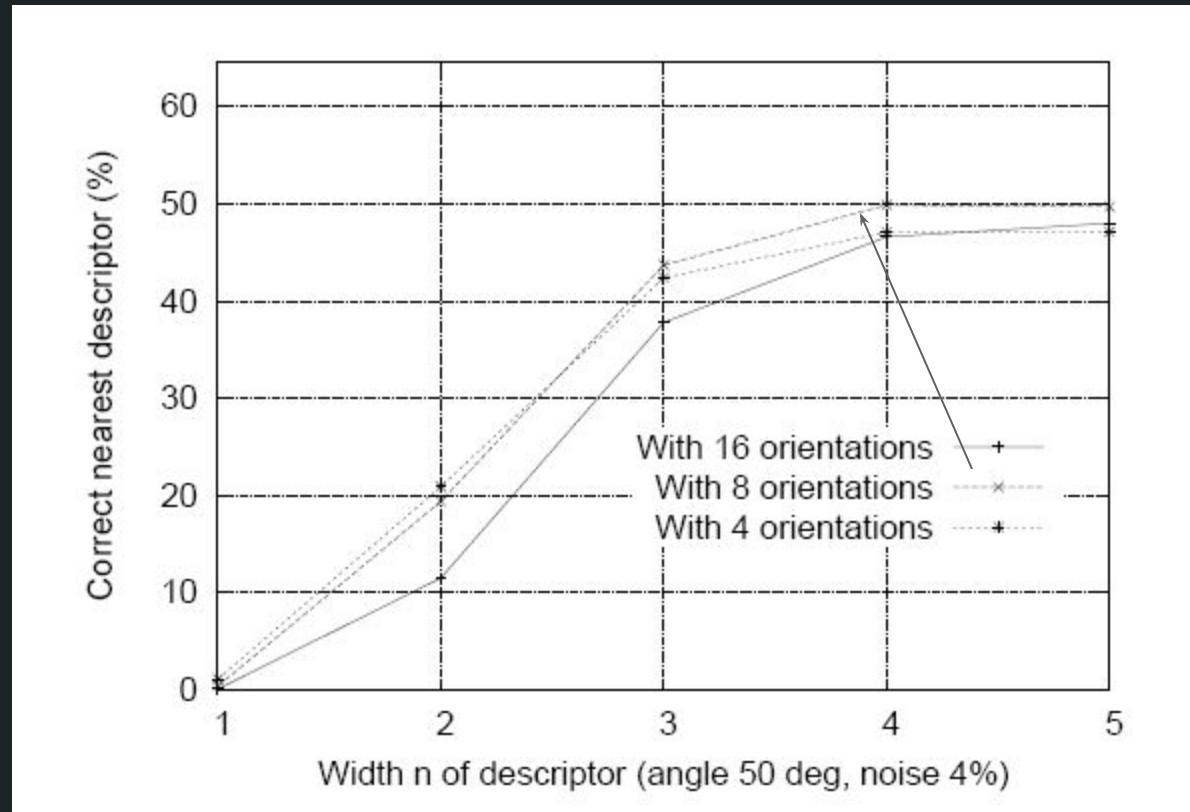
# SIFT Repeatability



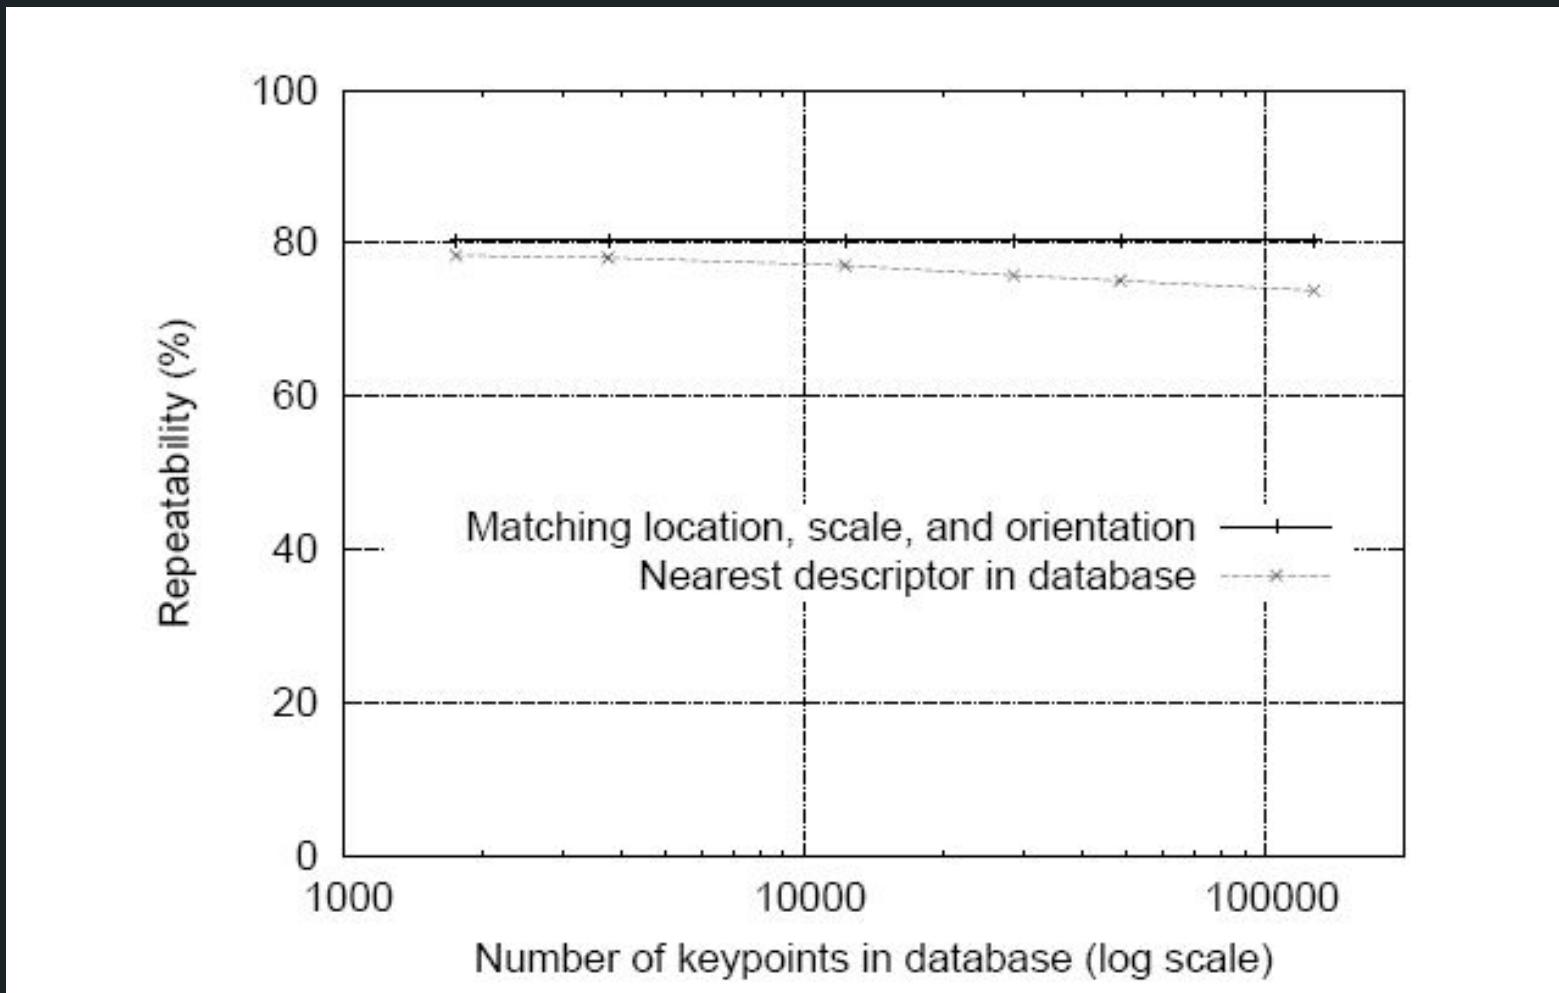
# SIFT Repeatability



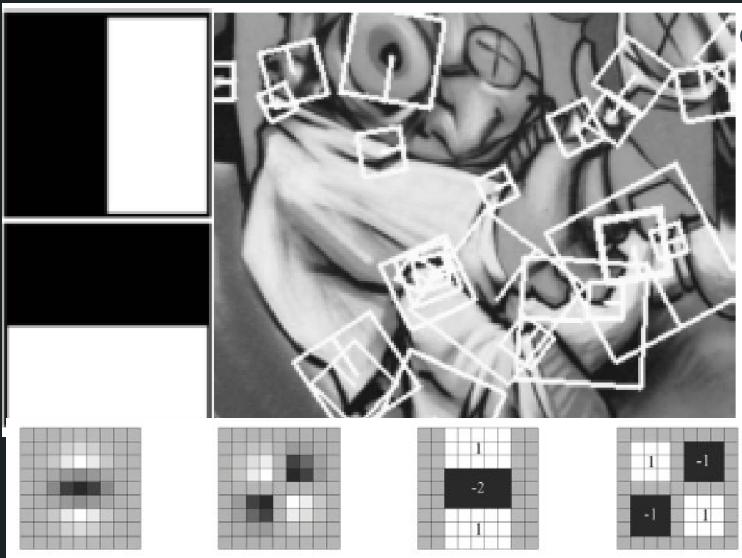
# SIFT Repeatability



# SIFT Repeatability

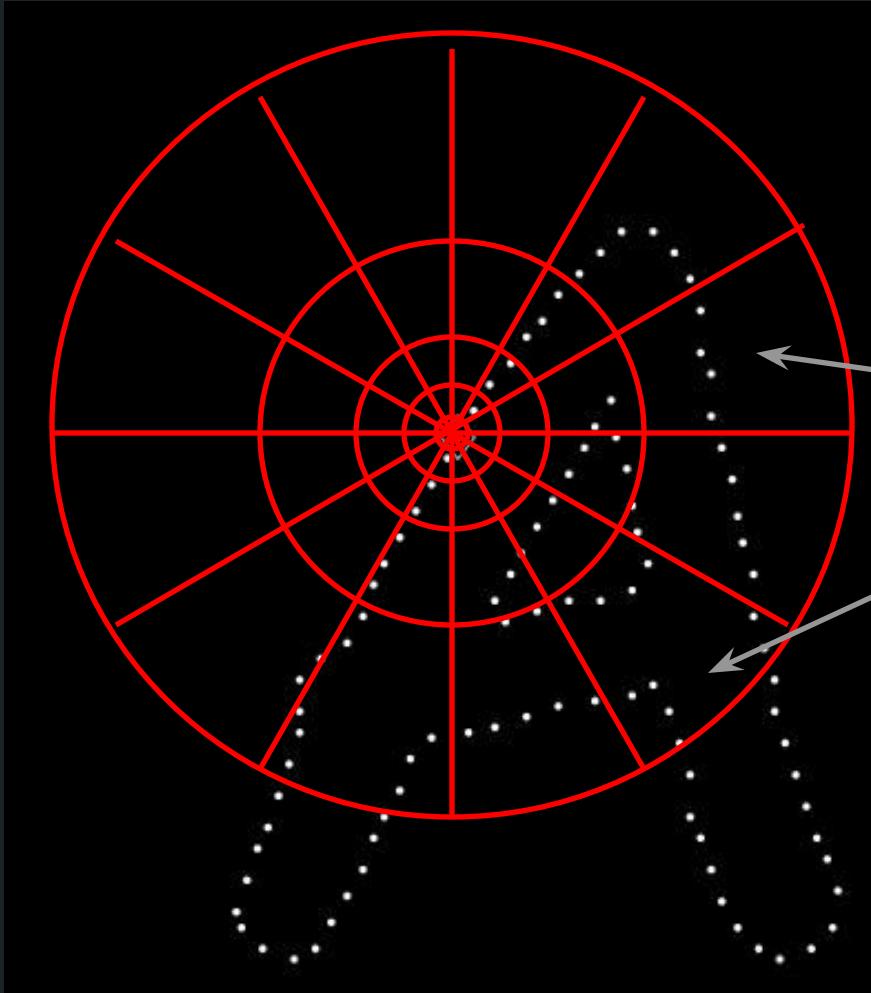


# Local Descriptors: SURF



- **Fast approximation of SIFT idea**
  - Efficient computation by 2D box filters & integral images  
⇒ 6 times faster than SIFT
  - Equivalent quality for object identification

## Local Descriptors: Shape Context



Count the number of points inside each bin, e.g.:

Count = 4

⋮

Count = 10

Log-polar binning: more precision for nearby points, more flexibility for farther points.

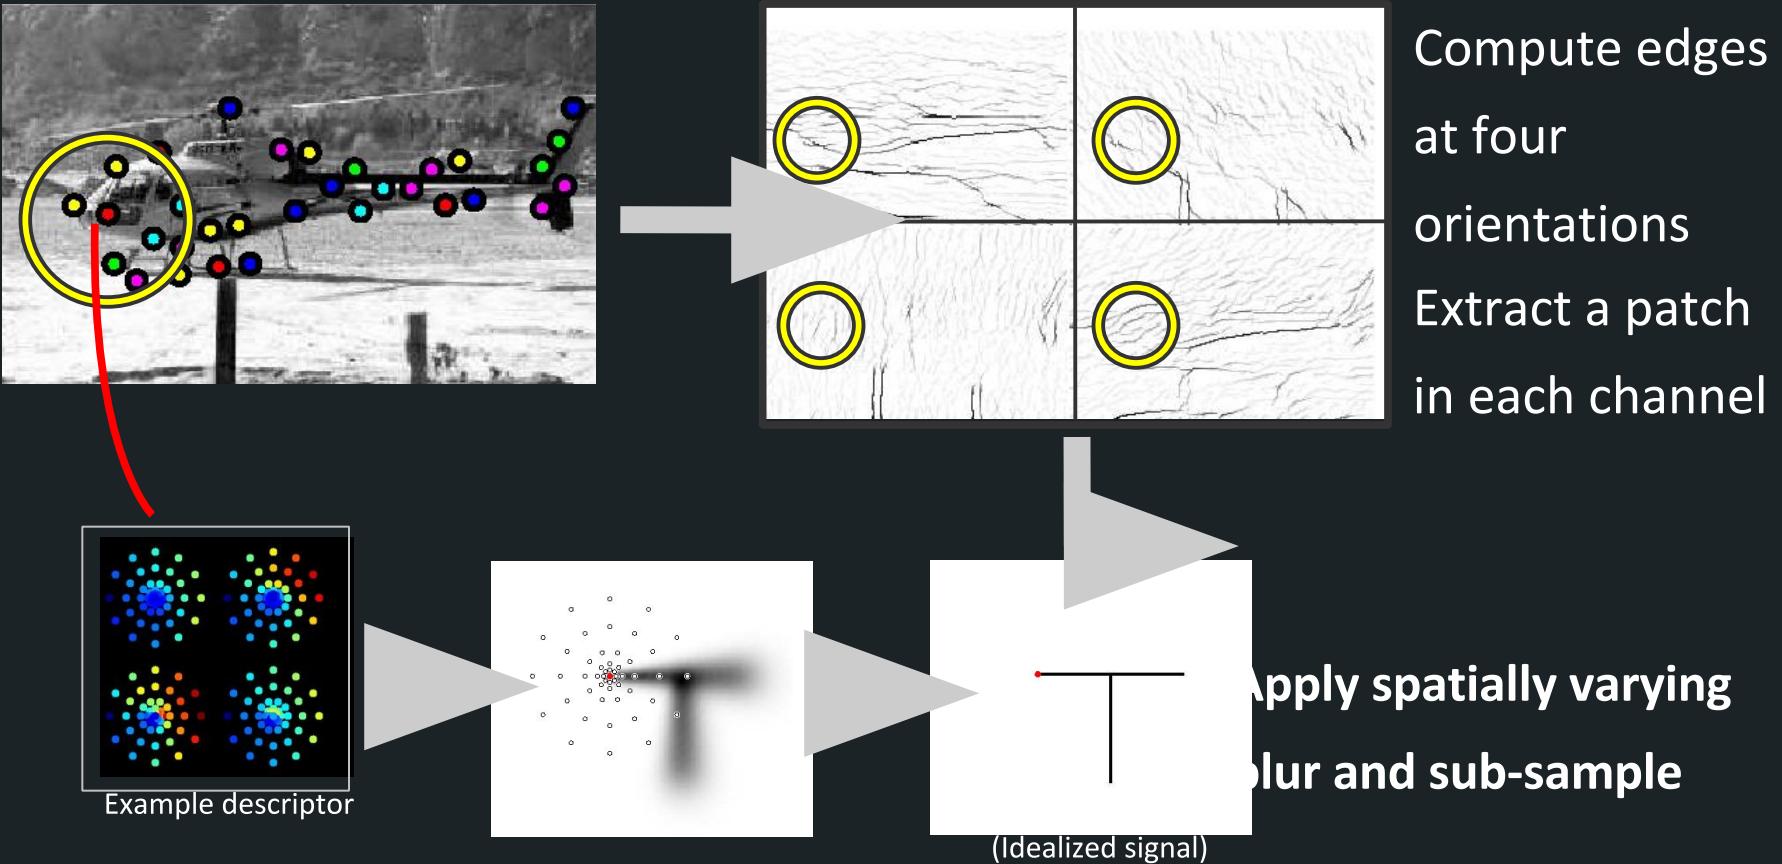
# SURF(Speeded-Up Robust Features)

Similar to SIFT but much faster



More like a blob detector

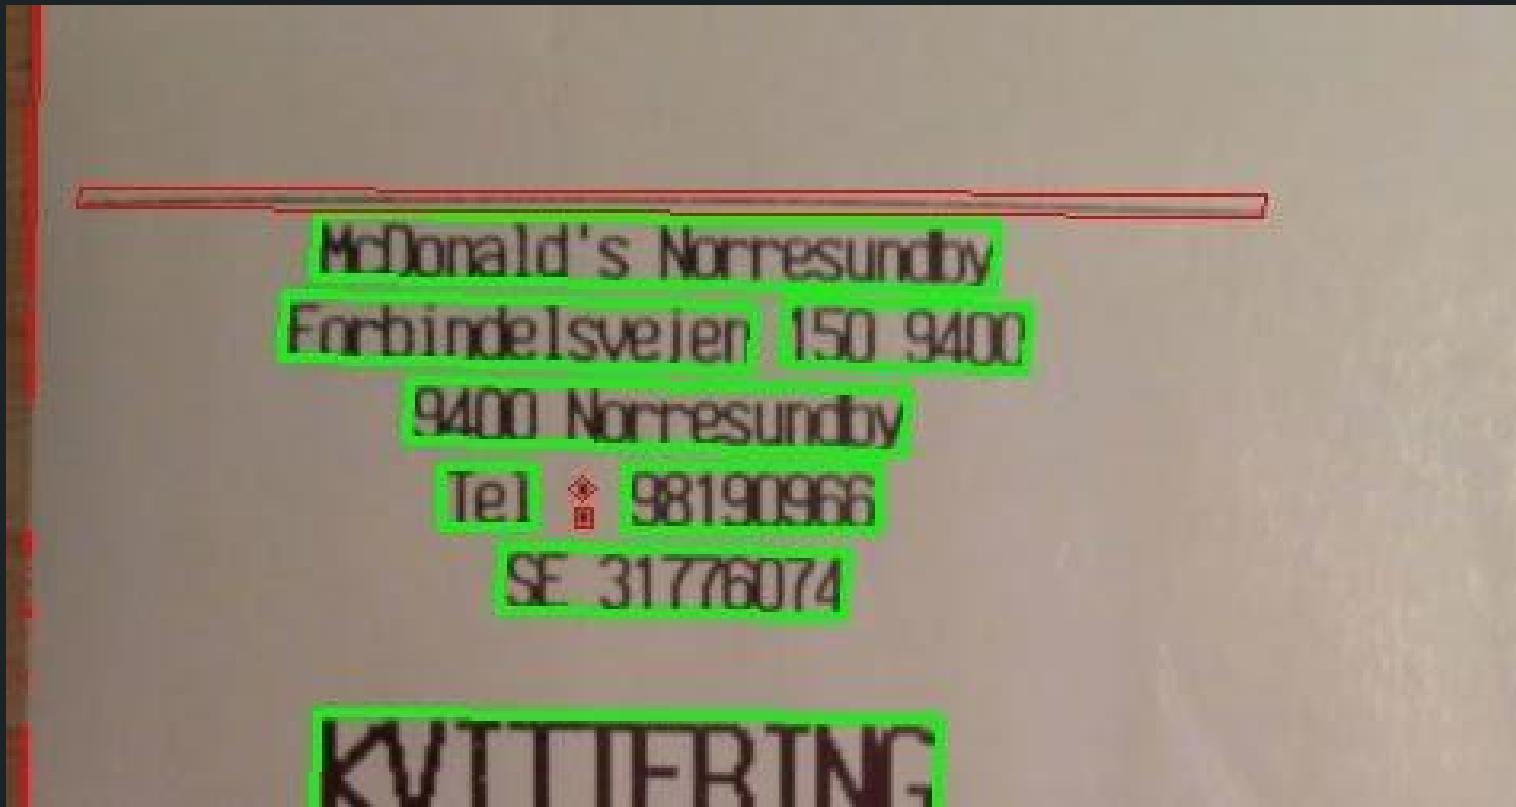
# Local Descriptors: Geometric Blur



# OTHER FEATURE DESCRIPTORS

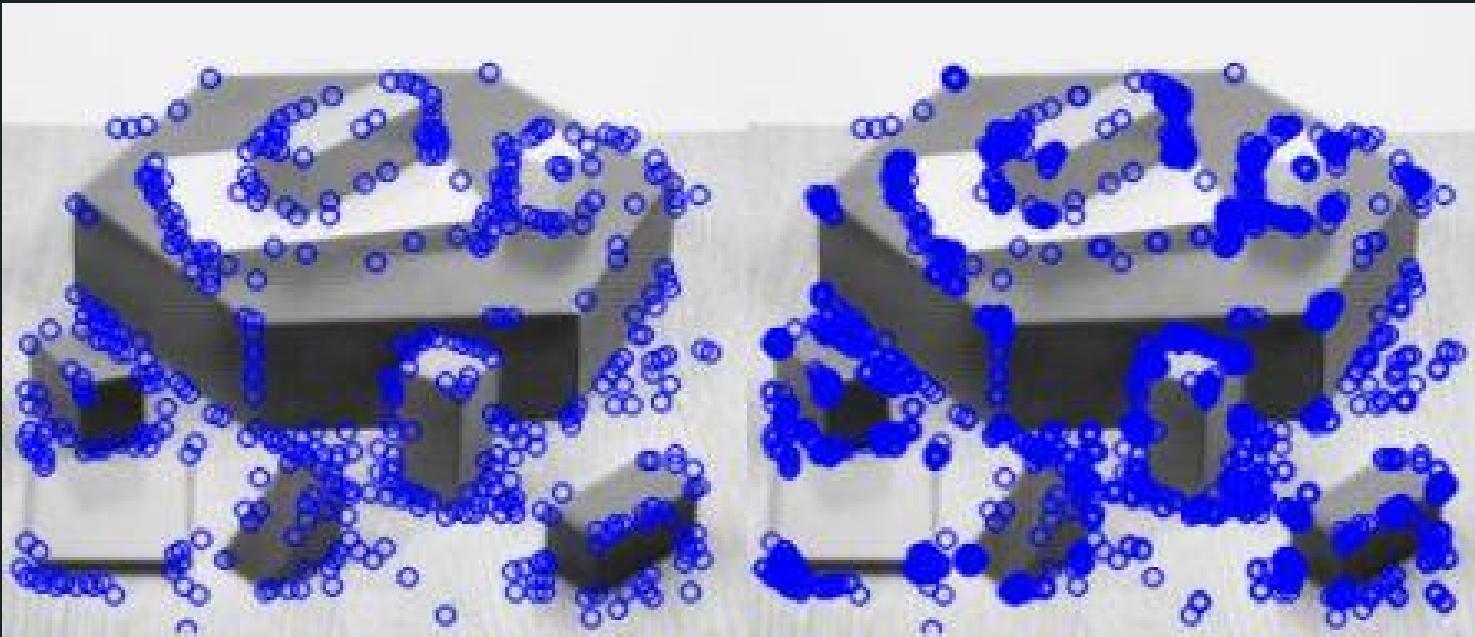
# MSER (Maximally Stable Extremal Regions)

## Blob Detection



# FAST(Features from Accelerated Segment Test)

Machine learning as a corner detector



First image shows FAST with nonmaxSuppression and second one without nonmaxSuppression

# Choosing a detector

- What do you want it for?
  - Precise localization in x-y: Harris
  - Good localization in scale: Difference of Gaussian
  - Flexible region shape: MSER
- Best choice often application dependent
  - Harris-/Hessian-Laplace/DoG work well for many natural categories
  - MSER works well for buildings and printed things
- Why choose?
  - Get more points with more detectors
- There have been extensive evaluations/comparisons
  - All detectors/descriptors shown here work well

# Comparison of Keypoint Detectors

Table 7.1 Overview of feature detectors.

Feature Detector	Corner	Blob	Region	Rotation invariant	Scale invariant	Affine invariant	Localization			
							Repeatability	accuracy	Robustness	Efficiency
Harris	✓			✓			+++	+++	+++	++
Hessian		✓		✓			++	++	++	+
SUSAN	✓			✓			++	++	++	+++
Harris-Laplace	✓	(✓)		✓	✓		+++	+++	++	+
Hessian-Laplace	(✓)	✓		✓	✓		+++	+++	+++	+
DoG	(✓)	✓		✓	✓		++	++	++	++
SURF	(✓)	✓		✓	✓		++	++	++	+++
Harris-Affine	✓	(✓)		✓	✓	✓	+++	+++	++	++
Hessian-Affine	(✓)	✓		✓	✓	✓	+++	+++	+++	++
Salient Regions	(✓)	✓		✓	✓	(✓)	+	+	++	+
Edge-based	✓			✓	✓	✓	+++	+++	+	+
MSER		✓		✓	✓	✓	+++	+++	++	+++
Intensity-based		✓		✓	✓	✓	++	++	++	++
Superpixels		✓		✓	(✓)	(✓)	+	+	+	+

# Choosing a descriptor

- Again, need not stick to one
- For object instance recognition or stitching, SIFT or variant is a good choice

# Things to remember

- Keypoint detection: repeatable and distinctive
  - Corners, blobs, stable regions
  - Harris, DoG



- Descriptors: robust and selective
  - spatial histograms of orientation
  - SIFT

